

Data Science

#HW4

Image Classification with Neural Network
산업공학과 201911532 송용재 (2023.10)

Table of Contents



Structure

Fully Connected Only

Convolution Neural Network

Building Block Method, ResNet50

Pytorch, torchvision

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix

import torch
import torchvision.models
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data.dataloader import DataLoader
from torchvision.datasets import ImageFolder
from torch.utils.data import random_split
import torchvision.transforms as transforms
from tqdm import tqdm
import torch.optim as optim

from torchvision.models import resnet50, ResNet50_Weights
import ssl
ssl._create_default_https_context = ssl._create_unverified_context
```

```
[14]    dir(torchvision.models)
[14]    ✓ 0.0s
...
'...',
'optical_flow',
'quantization',
'regnet',
'regnet_x_16gf',
'regnet_x_1_6gf',
'regnet_x_32gf',
'regnet_x_3_2gf',
'regnet_x_400mf',
'regnet_x_800mf',
'regnet_x_8gf',
'regnet_y_128gf',
'regnet_y_16gf',
'regnet_y_1_6gf',
'regnet_y_32gf',
'regnet_y_3_2gf',
'regnet_y_400mf',
'regnet_y_800mf',
'regnet_y_8gf',
'resnet',
'resnet101',
'resnet152',
'resnet18',
'resnet34',
'resnet50',
'resnext101_32x8d',
'resnext101_64x4d',
'resnext50_32x4d',
'segmentation',
'shufflenet_v2_x0_5',
'shufflenet_v2_x1_0'
```

개발환경



```
device = torch.device('mps') if torch.backends.mps.is_available() else "cpu"
print(device)
```

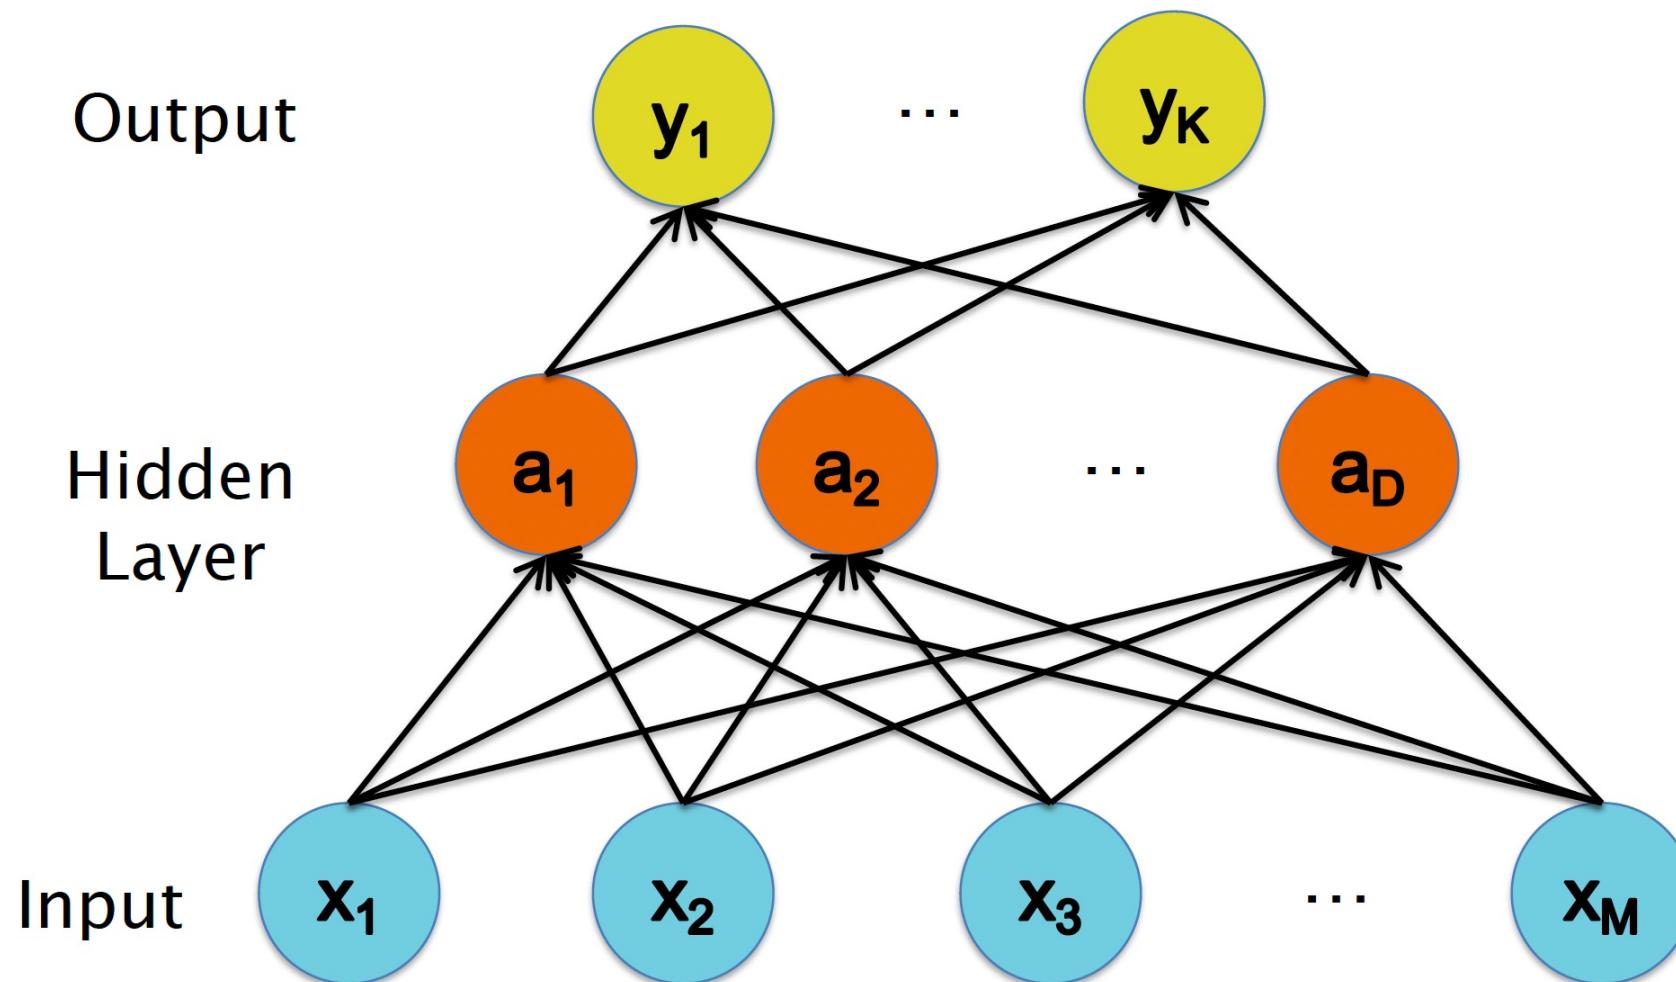
MPS BACKEND

... mps

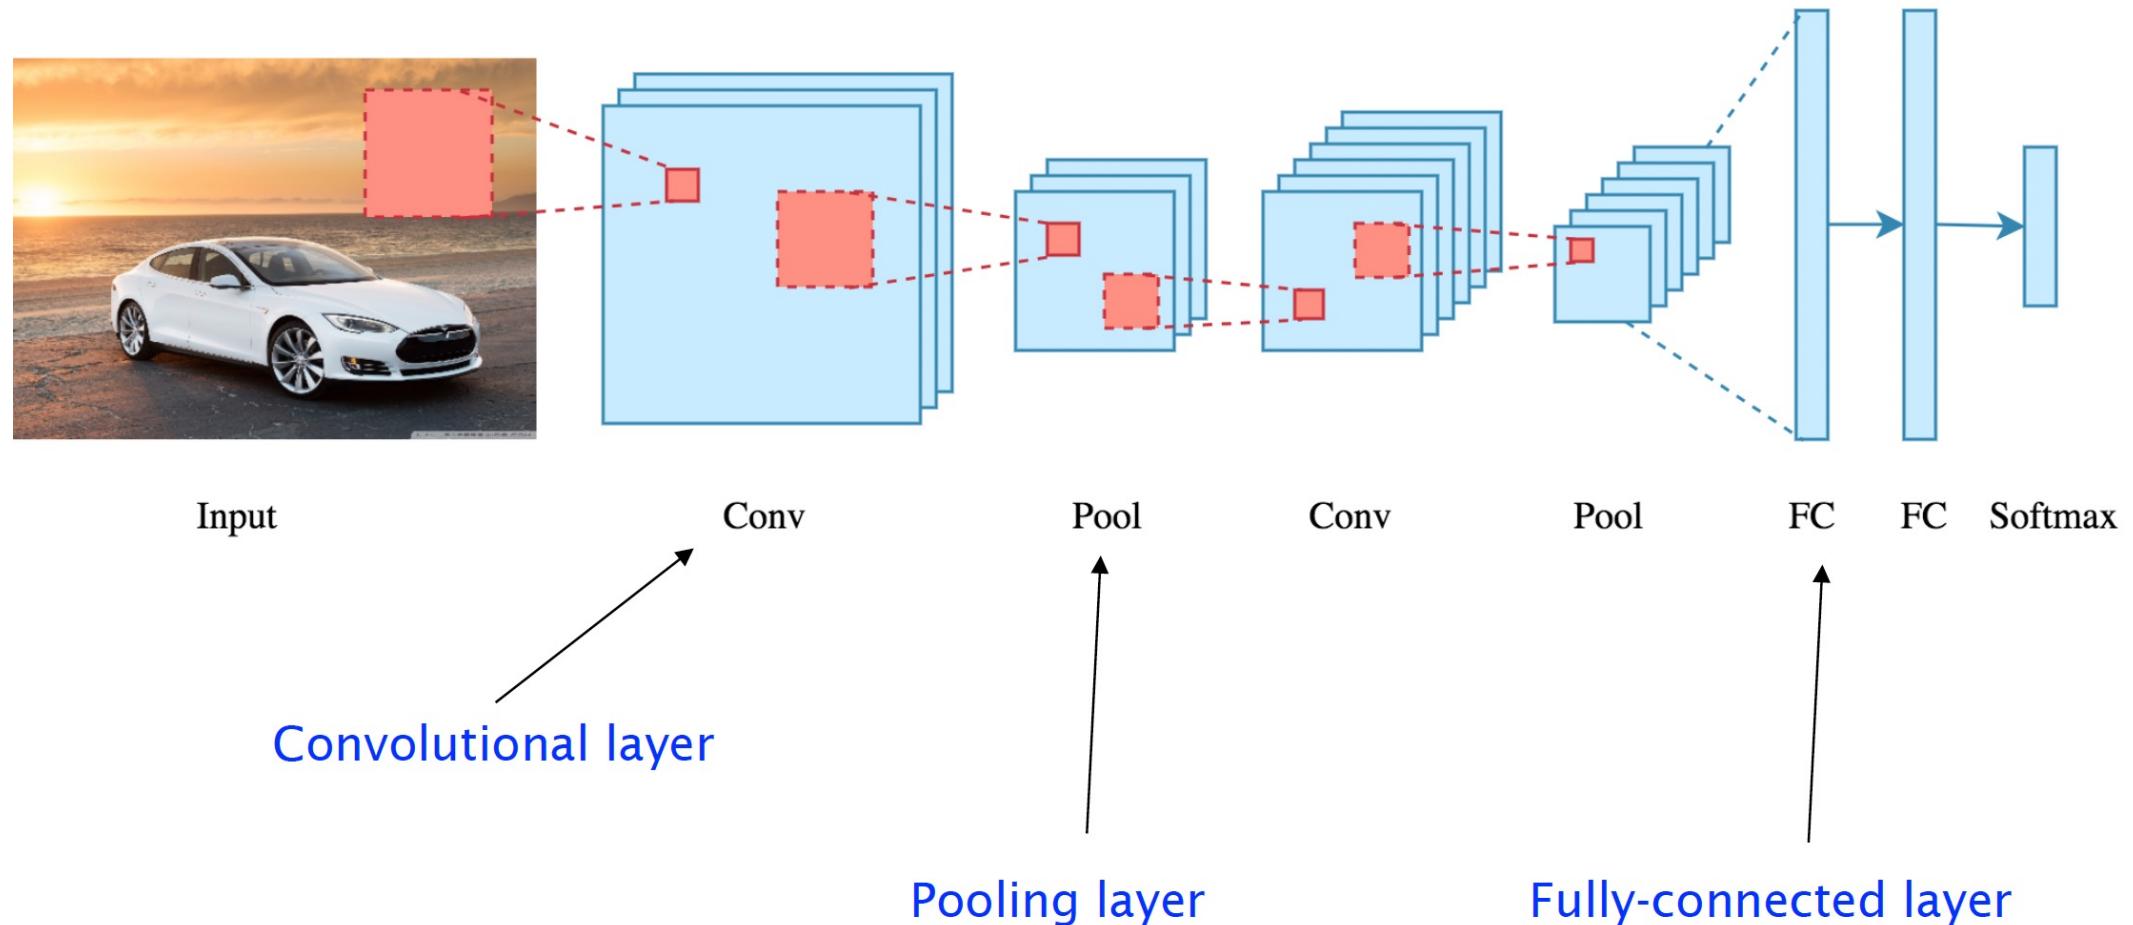
`mps` device enables high-performance training on GPU for MacOS devices with Metal programming framework. It introduces a new device to map Machine Learning computational graphs and primitives on highly efficient Metal Performance Shaders Graph framework and tuned kernels provided by Metal Performance Shaders framework respectively.

The new MPS backend extends the PyTorch ecosystem and provides existing scripts capabilities to setup and run operations on GPU.

Fully Connected Only



Convolution Neural Network



ResNet50

The ResNets developed in [1] are *modularized* architectures that stack building blocks of the same connecting shape. In this paper we call these blocks “*Residual*

Units”. The original Residual Unit in [1] performs the following computation:

$$\mathbf{y}_l = h(\mathbf{x}_l) + \mathcal{F}(\mathbf{x}_l, \mathcal{W}_l), \quad (1)$$

$$\mathbf{x}_{l+1} = f(\mathbf{y}_l). \quad (2)$$

Here \mathbf{x}_l is the input feature to the l -th Residual Unit. $\mathcal{W}_l = \{\mathbf{W}_{l,k} | 1 \leq k \leq K\}$ is a set of weights (and biases) associated with the l -th Residual Unit, and K is the number of layers in a Residual Unit (K is 2 or 3 in [1]). \mathcal{F} denotes the residual function, *e.g.*, a stack of two 3×3 convolutional layers in [1]. The function f is the operation after element-wise addition, and in [1] f is ReLU. The function h is set as an identity mapping: $h(\mathbf{x}_l) = \mathbf{x}_l$.¹

If f is also an identity mapping: $\mathbf{x}_{l+1} \equiv \mathbf{y}_l$, we can put Eqn.(2) into Eqn.(1) and obtain:

$$\mathbf{x}_{l+1} = \mathbf{x}_l + \mathcal{F}(\mathbf{x}_l, \mathcal{W}_l). \quad (3)$$

Recursively ($\mathbf{x}_{l+2} = \mathbf{x}_{l+1} + \mathcal{F}(\mathbf{x}_{l+1}, \mathcal{W}_{l+1}) = \mathbf{x}_l + \mathcal{F}(\mathbf{x}_l, \mathcal{W}_l) + \mathcal{F}(\mathbf{x}_{l+1}, \mathcal{W}_{l+1})$, etc.) we will have:

$$\mathbf{x}_L = \mathbf{x}_l + \sum_{i=l}^{L-1} \mathcal{F}(\mathbf{x}_i, \mathcal{W}_i), \quad (4)$$

ResNet50

Transfer Learning에 매우 효과적으로 사용될 수 있는 구조,

Residual Learning을 기반으로 Skip Connection을 도입하여 Gradient Vanishing 문제 해결의 초석 마련

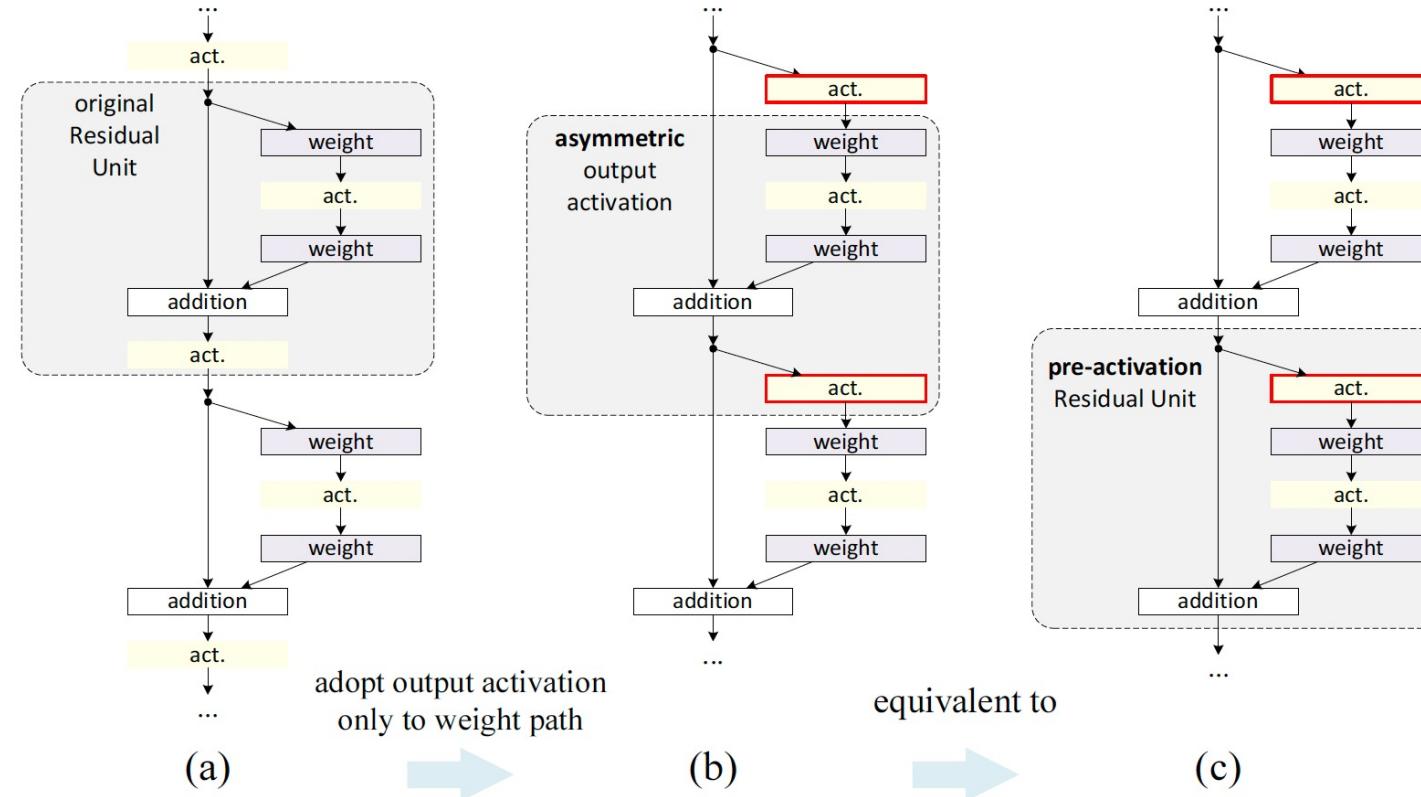


Figure 5. Using asymmetric after-addition activation is equivalent to constructing a pre-activation Residual Unit.

main 함수

```
if __name__ == '__main__':
    device = torch.device('mps') if torch.backends.mps.is_available() else "cpu"
    print(device)

    model1 = ImageClassificationModel_FC()
    model1.to(device)

    model2 = ImageClassificationModel_CNN()
    model2.to(device)

    resnet50(weights=ResNet50_Weights.DEFAULT)
    model3 = building_block_method(num_classes=3, fc_requires_grad=True)
    model3.to(device)

    dataloader = ImageDataLoader()

    trainer1 = ModelTrainer(model1, dataloader.train_loader, dataloader.val_loader, dataloader.test_loader, dataloader.classes, epochs=2)
    trainer1.train(device)
    trainer1.validate(device)
    trainer1.plot_loss()
    trainer1.evaluate(device)

    trainer2 = ModelTrainer(model2, dataloader.train_loader, dataloader.val_loader, dataloader.test_loader, dataloader.classes, epochs=2)
    trainer2.train(device)
    trainer2.validate(device)
    trainer2.plot_loss()
    trainer2.evaluate(device)

    trainer3 = ModelTrainer(model3, dataloader.train_loader, dataloader.val_loader, dataloader.test_loader, dataloader.classes, epochs=2)
    trainer3.train(device)
    trainer3.validate(device)
    trainer3.plot_loss()
    trainer3.evaluate(device)
```


Structure

Image Data Loader, 'train : validation : test = 7 : 1 : 2'

```
class ImageDataLoader:  
    def __init__(self):  
        transform = transforms.Compose([  
            transforms.Resize(300, 300),  
            transforms.ToTensor(),  
            transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])  
        ])  
  
        root_dir = 'ProjectImages'  
        full_train_dataset = ImageFolder(root=root_dir, transform=transform)  
        self.classes = full_train_dataset.classes  
  
        total_size = len(full_train_dataset)  
        train_size = int(total_size * 0.7)  
        val_size = int(total_size * 0.1)  
        test_size = total_size - train_size - val_size  
  
        self.train_data, self.val_data, self.test_data = random_split(full_train_dataset, [train_size, val_size, test_size])  
  
        self.train_loader = DataLoader(self.train_data, batch_size=32, shuffle=True, num_workers=2, pin_memory=True)  
        self.val_loader = DataLoader(self.val_data, batch_size=32, shuffle=True, num_workers=2, pin_memory=True)  
        self.test_loader = DataLoader(self.test_data, batch_size=32, shuffle=True, num_workers=2, pin_memory=True)
```

이미지 정규화

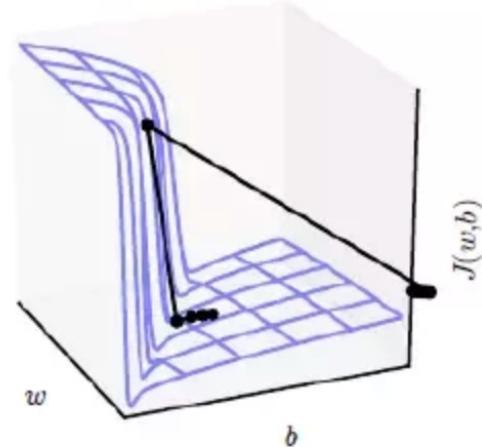
파일 호출 및 레이블 분류

Structure

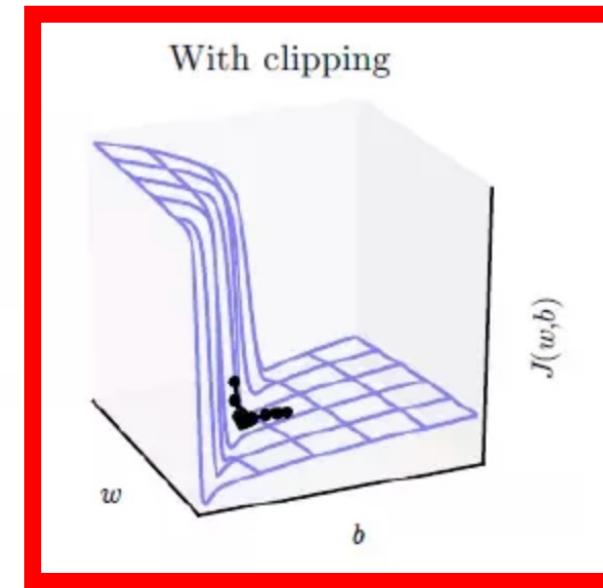
Class. Model Trainer

```
class ModelTrainer:  
    def __init__(self, model, train_loader, val_loader, test_loader, classes, epochs=10, learning_rate=0.001):  
        self.model = model  
        self.train_loader = train_loader  
        self.val_loader = val_loader  
        self.test_loader = test_loader  
        self.classes = classes  
        self.criterion = nn.CrossEntropyLoss()  
        self.optimizer = optim.Adam(model.parameters(), lr=learning_rate)  
  
        self.epochs = epochs  
  
        self.train_losses = []  
        self.val_losses = []
```

Without clipping



With clipping



Class_Model_Trainer

Model training with Progress Bar (tqdm)

```
def train(self, device):
    self.model.to(device)
    for epoch in range(self.epochs):
        self.model.train()
        train_loss = 0
        with tqdm(total=len(self.train_loader), desc=f"Epoch {epoch + 1}/{self.epochs}", unit="batch", leave=False) as pbar:
            for features, labels in self.train_loader:
                features, labels = features.to(device), labels.to(device)
                outputs = self.model(features)
                loss = self.criterion(outputs, labels)
                train_loss += loss.item()

                self.optimizer.zero_grad()
                loss.backward()
                self.optimizer.step()

                pbar.set_postfix({'train_loss': f'{loss.item():.4f}'})
                pbar.update(1)

        train_loss /= len(self.train_loader)
        val_loss = self.validate(device)

        self.train_losses.append(train_loss)
        self.val_losses.append(val_loss)

        print(f"Epoch {epoch + 1}/{self.epochs} - Train Loss: {train_loss:.4f}, Validation Loss: {val_loss:.4f}")
```

최적화 및 update 저장

Class_Model_Trainer

Validate & Plot loss

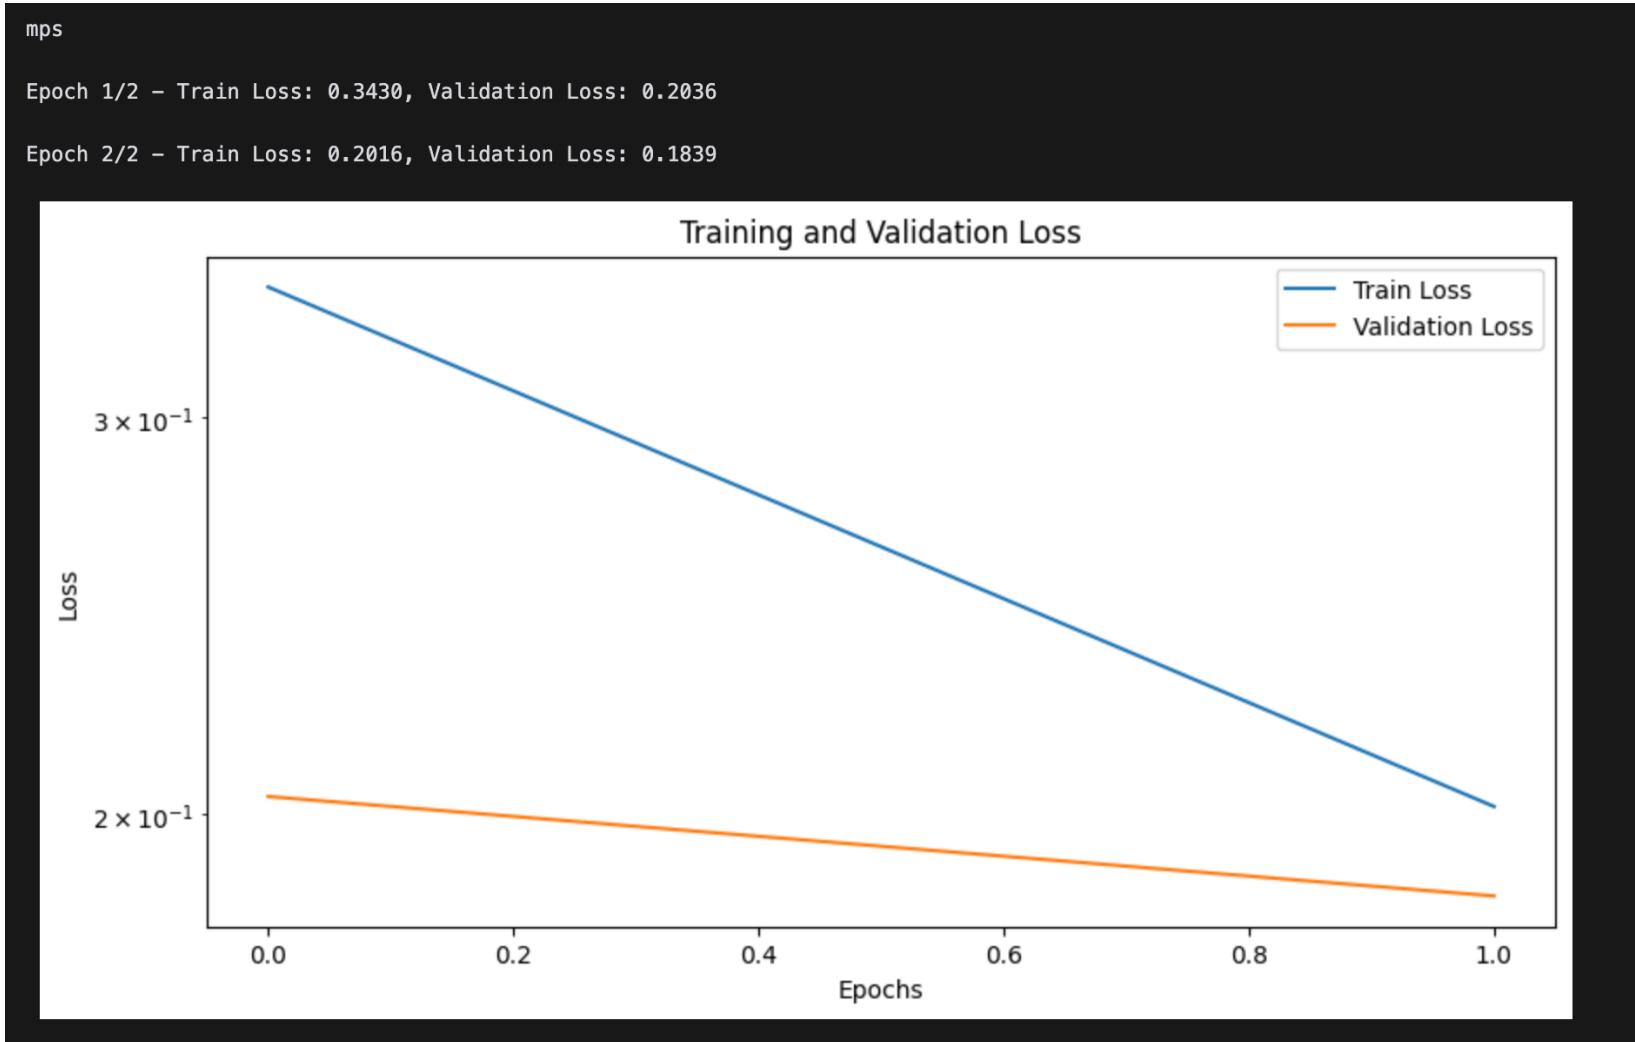
```
def validate(self, device):
    self.model.eval()
    val_loss = 0
    with torch.no_grad():
        for features, labels in self.val_loader:
            features, labels = features.to(device), labels.to(device)
            outputs = self.model(features)
            loss = self.criterion(outputs, labels)
            val_loss += loss.item()

    val_loss /= len(self.val_loader)
    return val_loss

def plot_loss(self):
    plt.figure(figsize=(10, 5))
    plt.plot(self.train_losses, label='Train Loss')
    plt.plot(self.val_losses, label='Validation Loss')
    plt.title('Training and Validation Loss')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.yscale('log')
    plt.legend()
    plt.show()
```

Class_Model_Trainer

'Train Loss, Validation Loss & Plot Loss' Examples



Class_Model_Trainer

Evaluate

Output으로의 분류를 위한
Softmax 함수

$$S(x_i) = \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}}$$

```

def evaluate(self, device):
    self.model.eval()
    test_loss = 0
    all_labels = []
    all_preds = []

    total = 0
    correct = 0

    correct_predictions = []
    incorrect_predictions = []

    with torch.no_grad():
        for features, labels in self.test_loader:
            features, labels = features.to(device), labels.to(device)
            outputs = self.model(features.float())
            loss = self.criterion(outputs, labels)
            test_loss += loss.item()

            output_prob = F.softmax(outputs, dim=1).cpu().numpy()

            _, predicted = torch.max(outputs, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

            all_labels.extend(labels.cpu().numpy())
            all_preds.extend(predicted.cpu().numpy())

        for f, l, p, o, prob in zip(features, labels, predicted, outputs, output_prob):
            if l == p and len(correct_predictions) < 10:
                correct_predictions.append((f, l, p, o, prob))
            elif l != p and len(incorrect_predictions) < 10:
                incorrect_predictions.append((f, l, p, o, prob))
            if len(correct_predictions) >= 10 and len(incorrect_predictions) >= 10:
                break

    print("\n----- Correct Predictions -----")

```

Class_Model_Trainer

Evaluate, '맞은 예측 10개 & 틀린 예측 10개 + Accuracy & Confusion Matrix'

```

print("\n----- Correct Predictions -----")
for f, l, p, o, prob in correct_predictions:
    f_permute = f.permute(1, 2, 0).cpu().numpy()

    restored_image = (f_permute * [0.229, 0.224, 0.225]) + [0.485, 0.456, 0.406]
    restored_image = np.clip(restored_image, 0, 1)

    plt.imshow(restored_image, cmap='gray')
    plt.title(f'Predicted: {self.classes[p]}, Actual: {self.classes[l]}')
    plt.show()
    print(f"{o}-->{prob}: {self.classes[l]} -> {self.classes[p]}')

print("\n----- Incorrect Predictions -----")
for f, l, p, o, prob in incorrect_predictions:
    f_permute = f.permute(1, 2, 0).cpu().numpy()

    restored_image = (f_permute * [0.229, 0.224, 0.225]) + [0.485, 0.456, 0.406]
    restored_image = np.clip(restored_image, 0, 1)

    plt.imshow(restored_image, cmap='gray')
    plt.title(f'Predicted: {self.classes[p]}, Actual: {self.classes[l]}')
    plt.show()
    print(f"{o}-->{prob}: {self.classes[l]} -> {self.classes[p]}')

test_loss /= len(self.test_loader)
print(f"Accuracy = {correct/total: .4f}")

cm = confusion_matrix(all_labels, all_preds)
print(f"Test Loss: {test_loss:.4f}")
print("Confusion Matrix:")
print(cm)

```

정규화 된 이미지를
다시 원 상태로 복구

Fully Connected Only

nn.Module

Hidden Layer는 6개, hidden layer 대체로 128×128 형태, Dropout(50%)을 통해 Overfitting 방지

```
class ImageClassificationModel_FC(nn.Module):
    def __init__(self):
        super(ImageClassificationModel_FC, self).__init__()
        self.fc1 = nn.Linear(3 * 300 * 300, 256)
        self.fc2 = nn.Linear(256, 128)
        self.fc3 = nn.Linear(128, 128)
        self.fc4 = nn.Linear(128, 128)
        self.fc5 = nn.Linear(128, 128)
        self.fc6 = nn.Linear(128, 128)

        self.dropout = nn.Dropout(0.5)

        self.fcm = nn.Linear(128, 3)

    def forward(self, x):
        x = x.reshape(x.size(0), -1)
        x = torch.relu(self.fc1(x))
        x = torch.relu(self.fc2(x))
        x = torch.relu(self.fc3(x))
        x = torch.relu(self.fc4(x))
        x = torch.relu(self.fc5(x))
        x = torch.relu(self.fc6(x))

        x = self.dropout(x)

        x = self.fcm(x)

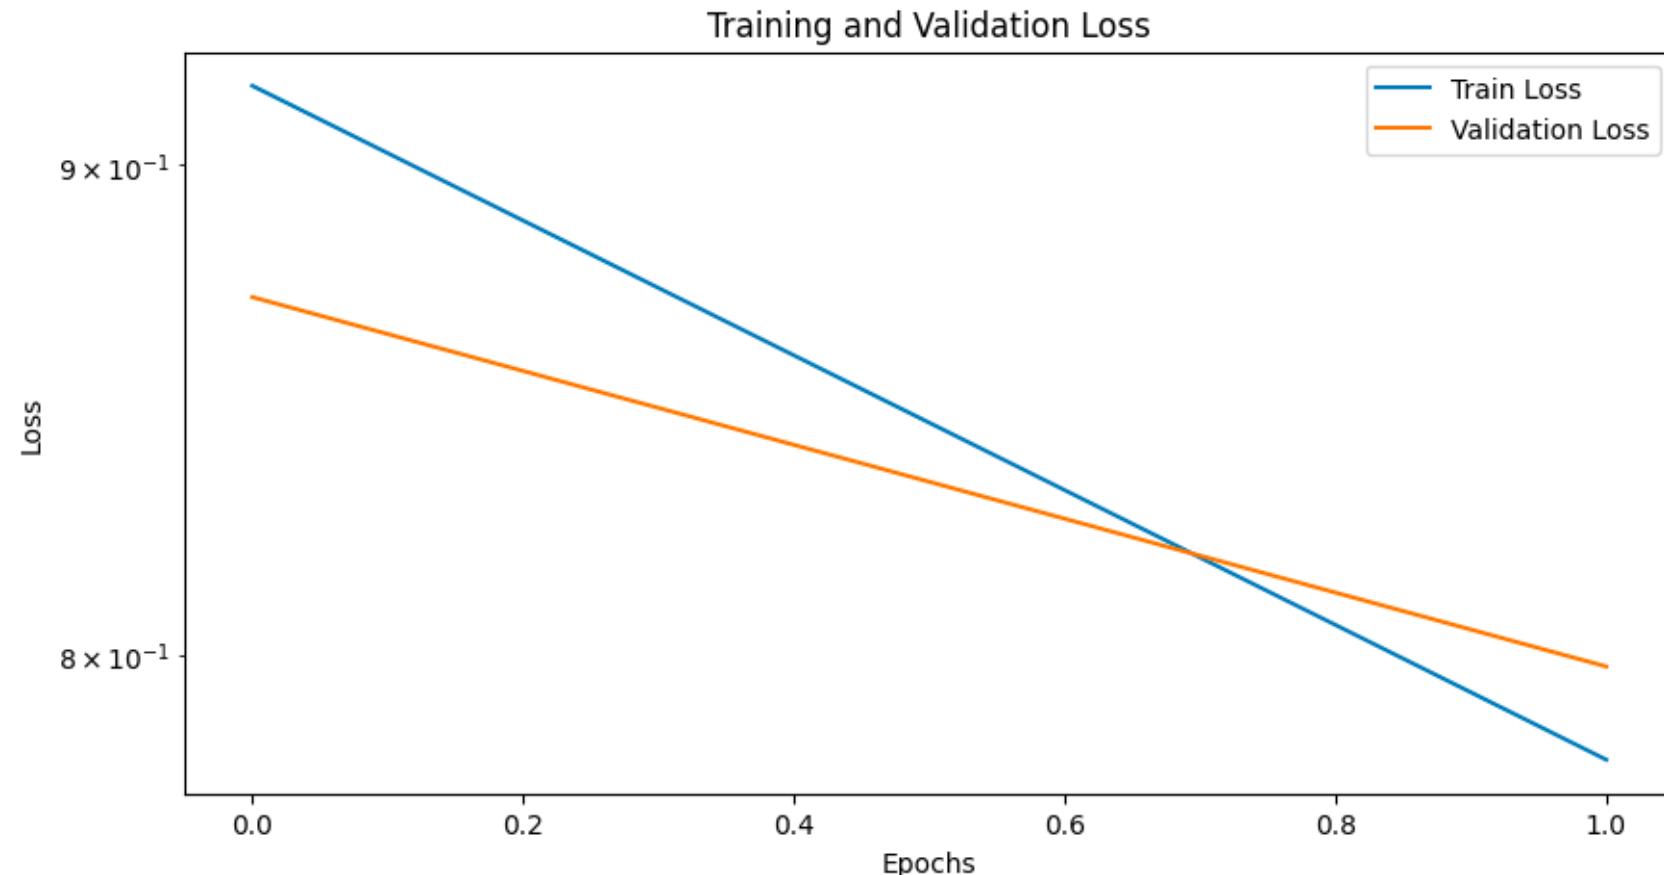
    return x
```

Fully Connected Only

Plot Loss

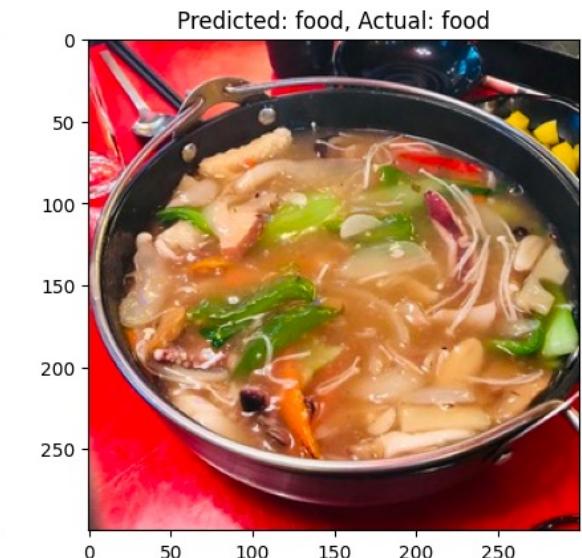
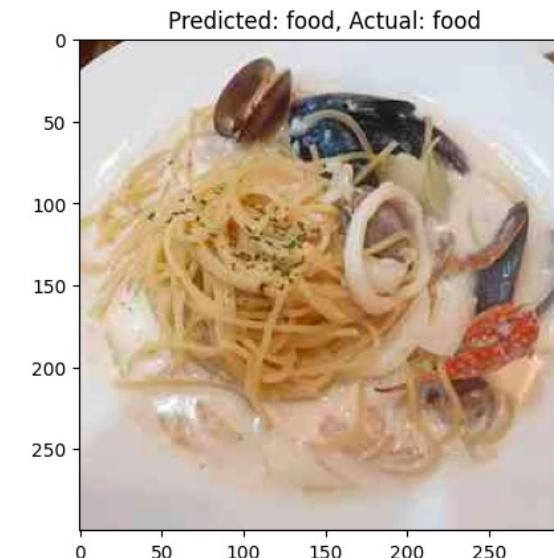
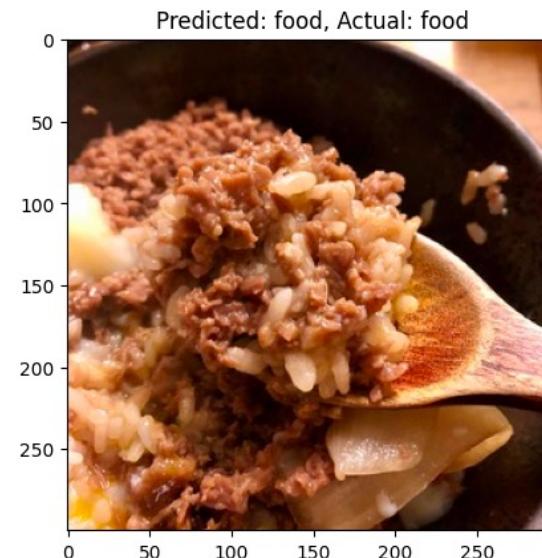
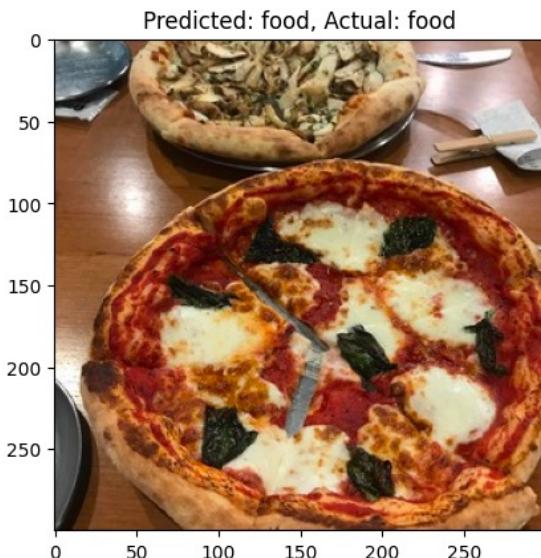
Epoch 1/2 – Train Loss: 0.9172, Validation Loss: 0.8718

Epoch 2/2 – Train Loss: 0.7801, Validation Loss: 0.7978



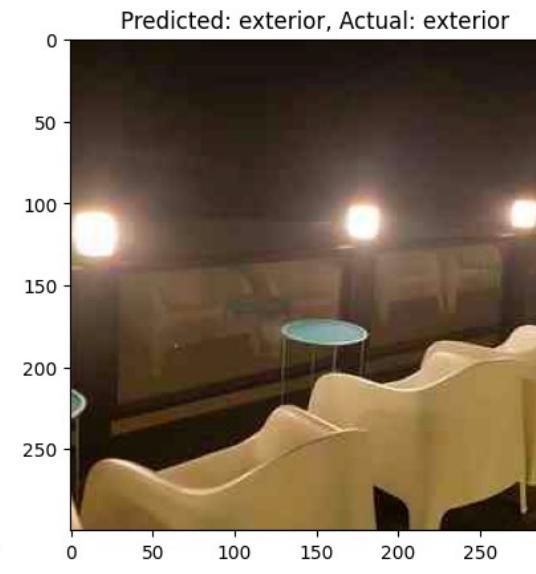
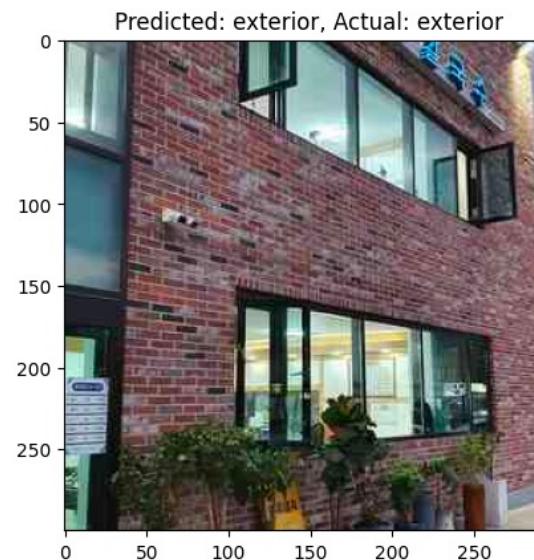
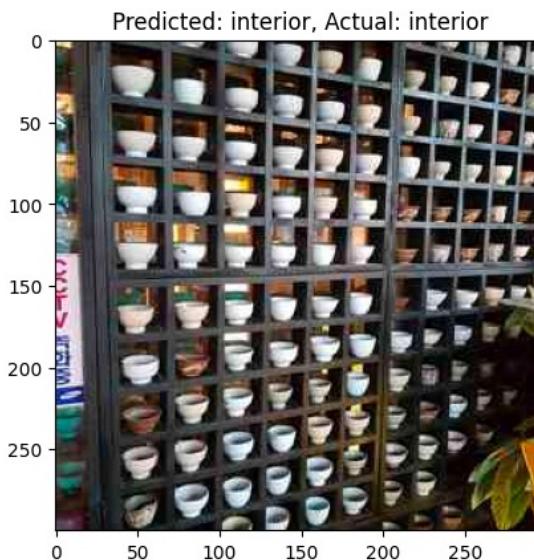
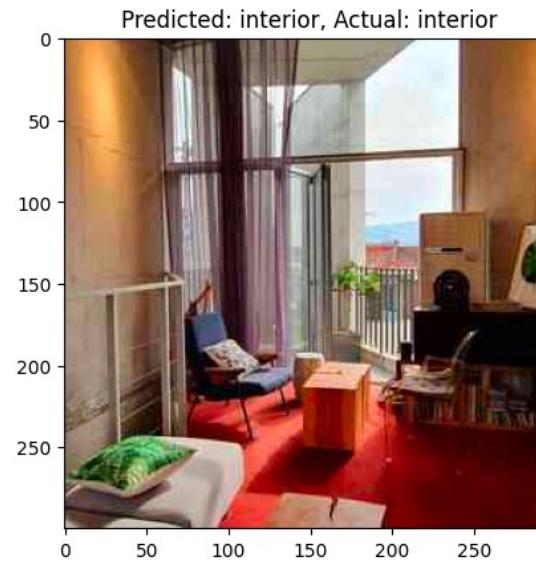
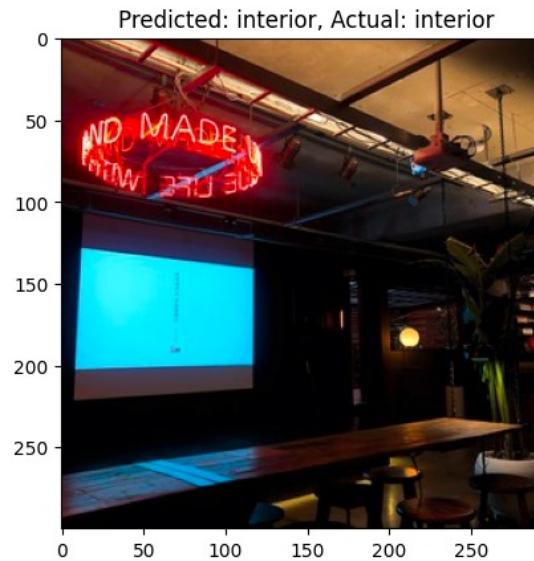
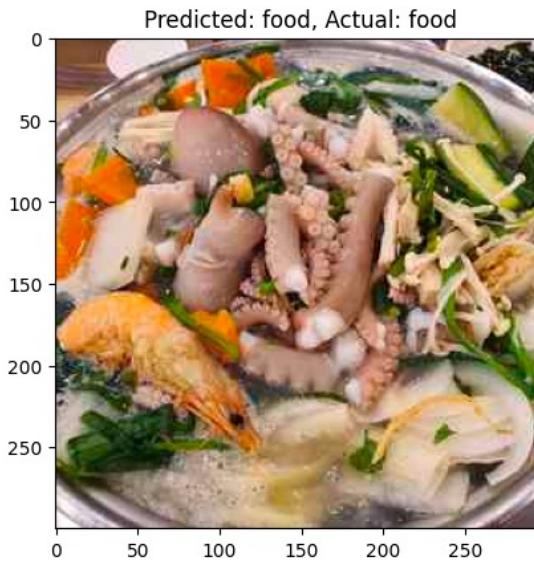
Fully Connected Only

----Correct Predictions---- 10개



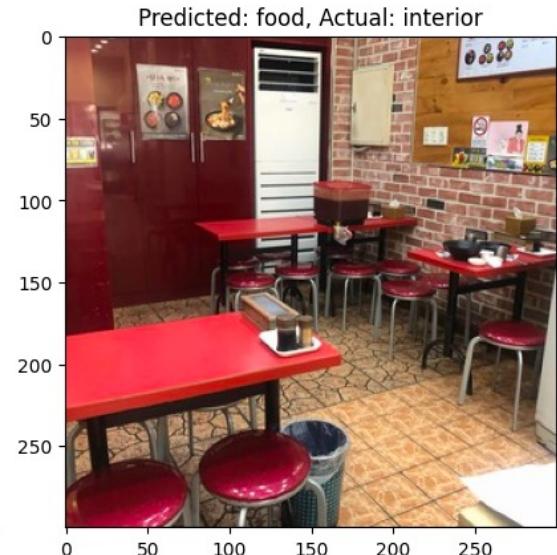
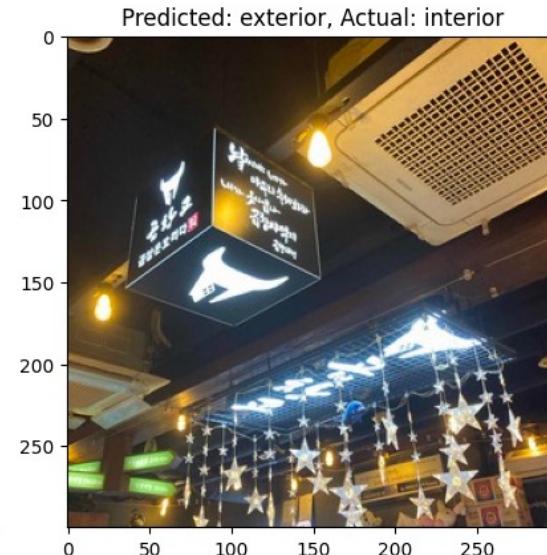
Fully Connected Only

----Correct Predictions----



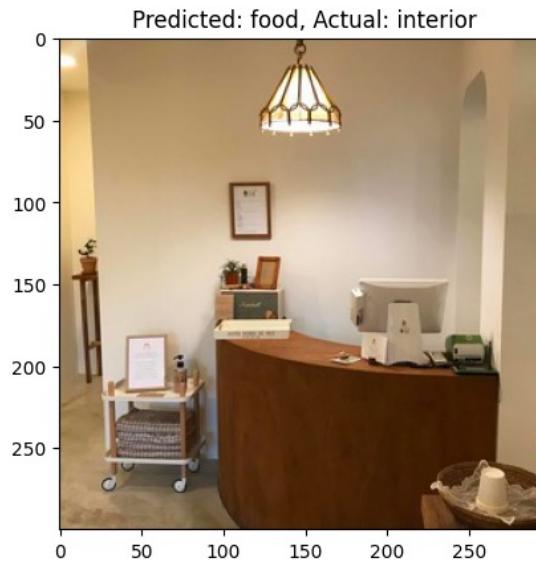
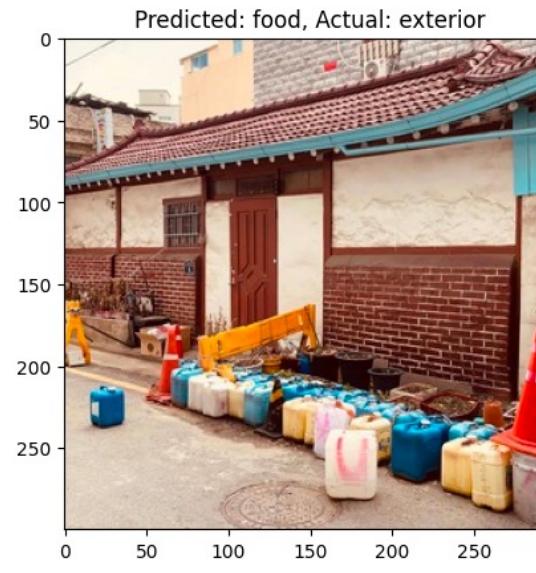
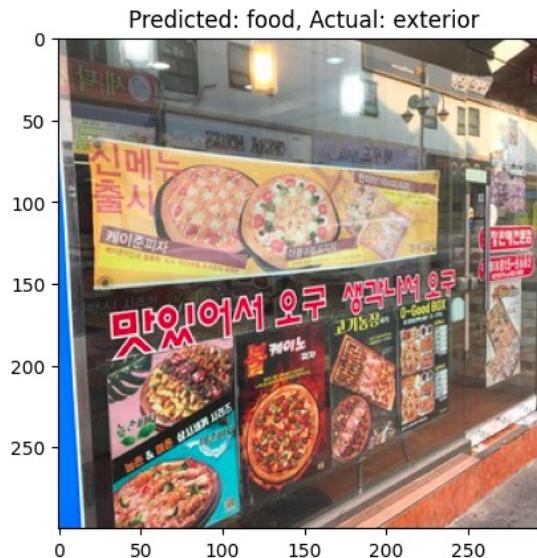
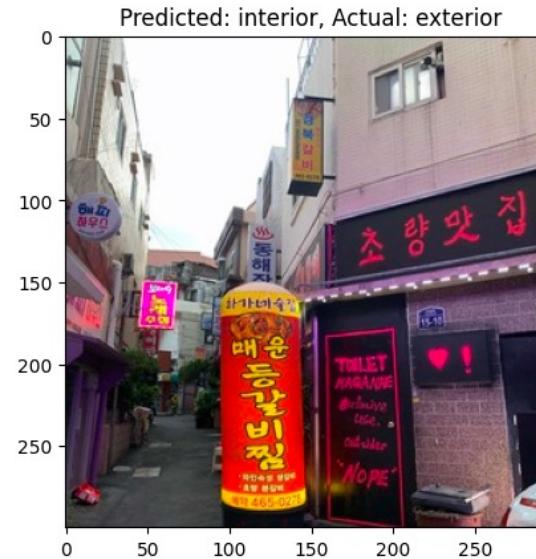
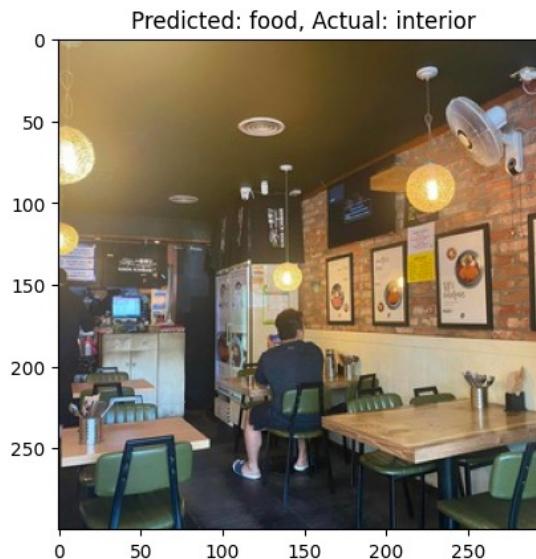
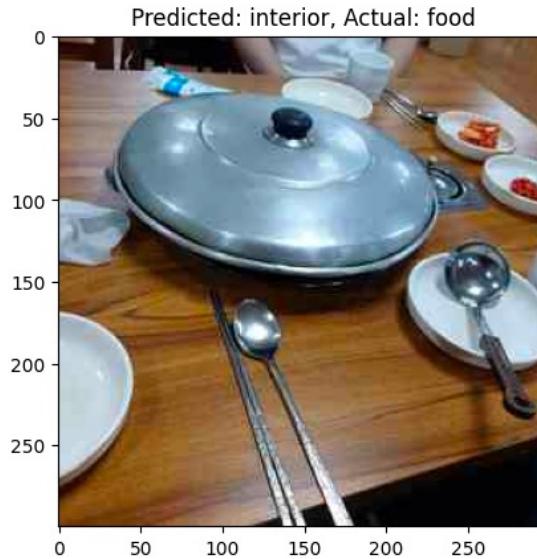
Fully Connected Only

----Incorrect Predictions---- 10개



Fully Connected Only

----Incorrect Predictions----



Accuracy & Confusion Matrix

Accuracy : 0.6367, 사용하기 많이 애매한 정확도

Accuracy = 0.6367

Test Loss: 0.8397

Confusion Matrix:

[[172 64 76]

 [42 513 88]

 [96 156 230]]

Convolution Neural Network

nn.Module

총 3개의 Convolution-Pooling 쌍, Stride=1 + Padding=1을 통해 kernel_size에 상응하는 input-output size matching

```

class ImageClassificationModel_CNN(nn.Module):
    def __init__(self):
        super(ImageClassificationModel_CNN, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, kernel_size=3, stride=1, padding=1)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1)
        self.conv3 = nn.Conv2d(64, 64, kernel_size=3, stride=1, padding=1)

        self.pool = nn.MaxPool2d(kernel_size=2, stride=2, padding = 0)
        self.dropout = nn.Dropout(0.5)

        self.fc1 = nn.Linear(64 * 37 * 37, 64)
        self.fc2 = nn.Linear(64, 32)
        self.fc3 = nn.Linear(32, 3)

    def forward(self, x):
        x = self.pool(torch.relu(self.conv1(x)))
        x = self.pool(torch.relu(self.conv2(x)))
        x = self.pool(torch.relu(self.conv3(x)))

        x = x.view(-1, 64 * 37 * 37)
        x = self.dropout(x)
        x = torch.relu(self.fc1(x))
        x = torch.relu(self.fc2(x))
        x = self.fc3(x)

    return x

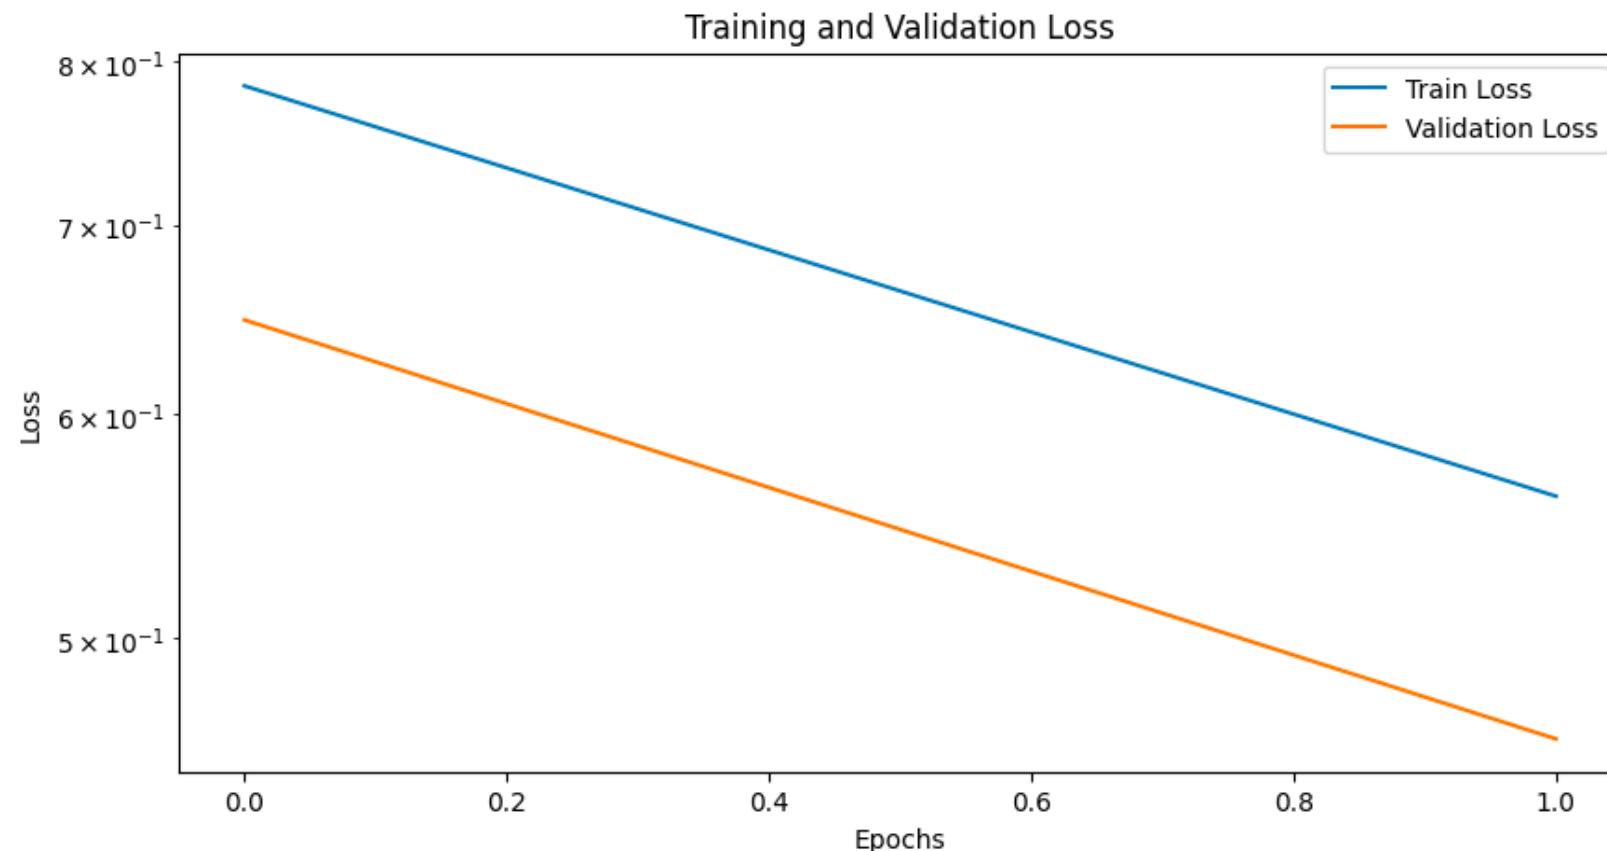
```

Dropout 이후 구성된
세 개의 FC

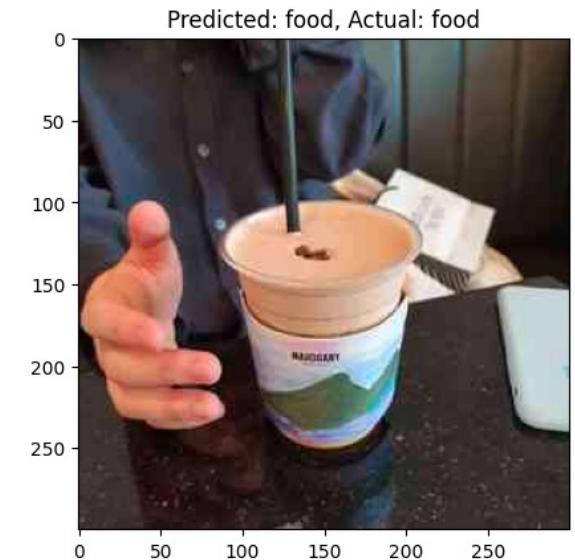
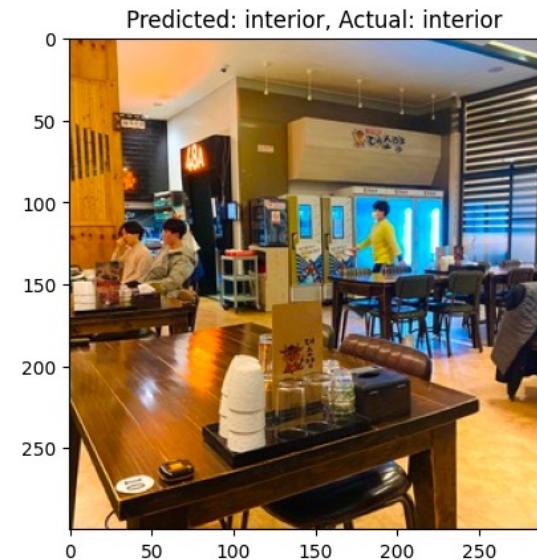
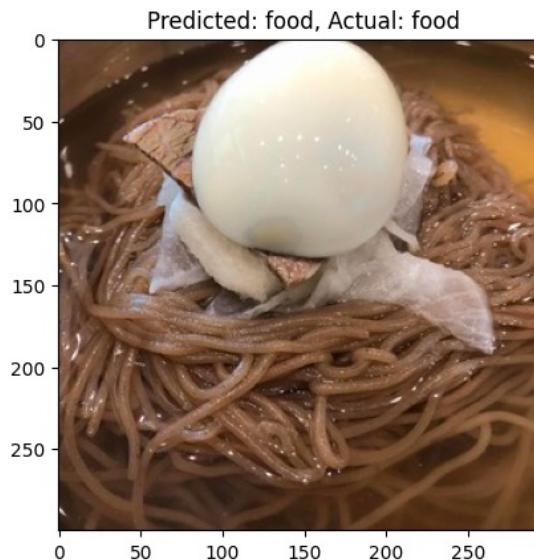
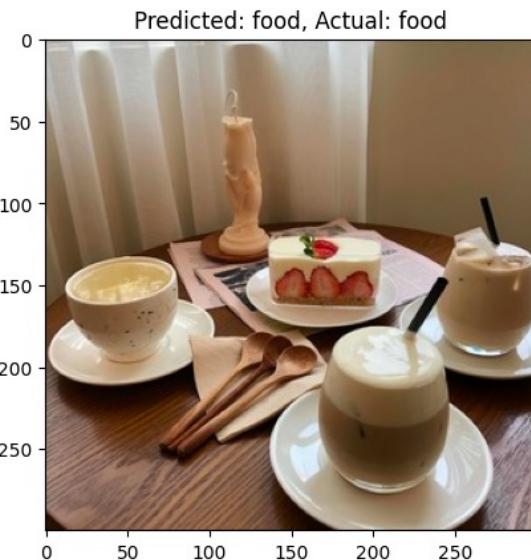
Plot Loss

Epoch 1/2 – Train Loss: 0.7843, Validation Loss: 0.6481

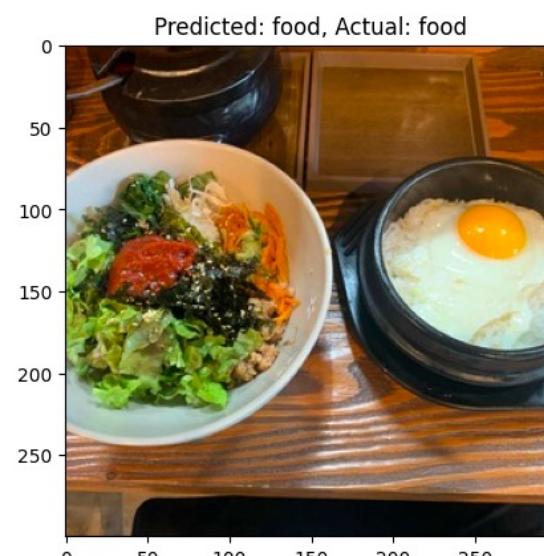
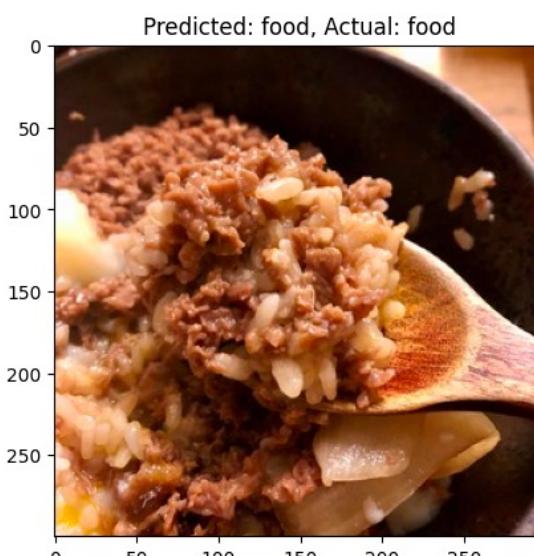
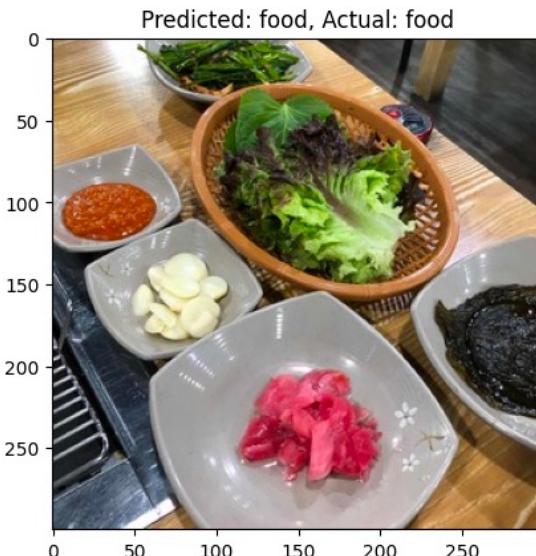
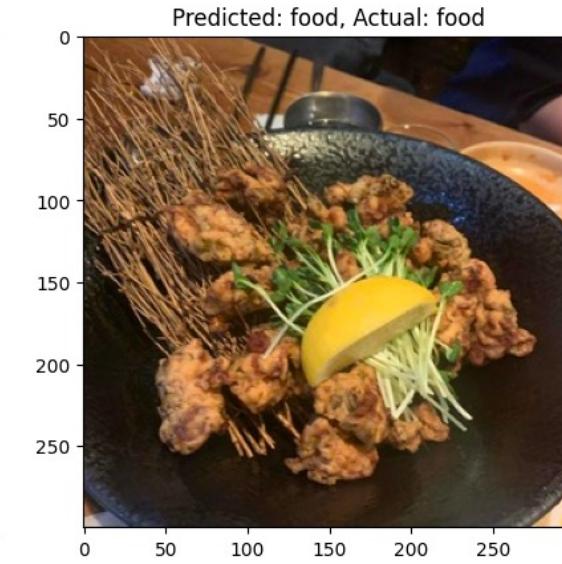
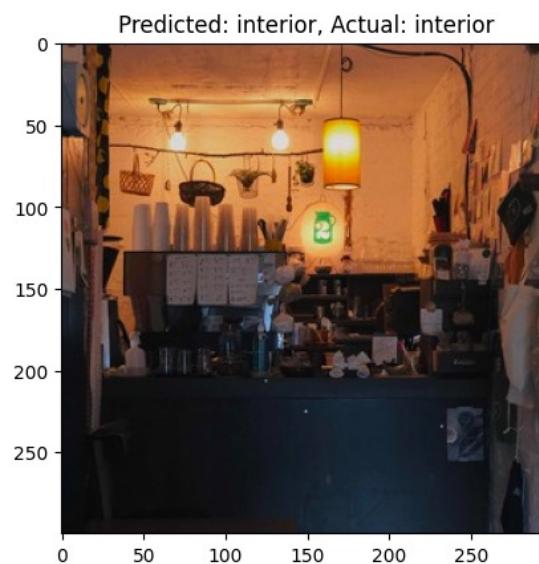
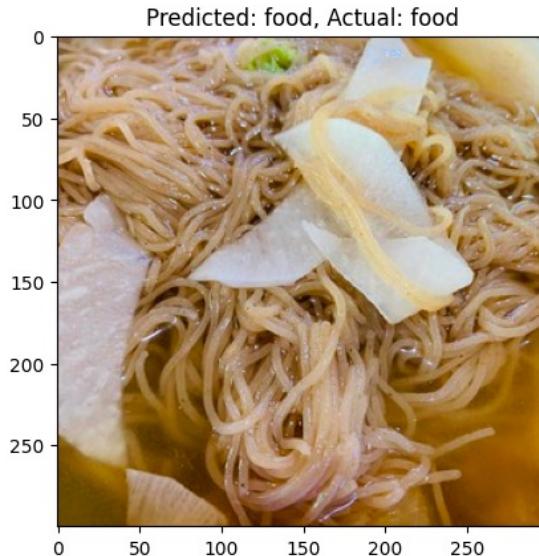
Epoch 2/2 – Train Loss: 0.5615, Validation Loss: 0.4608



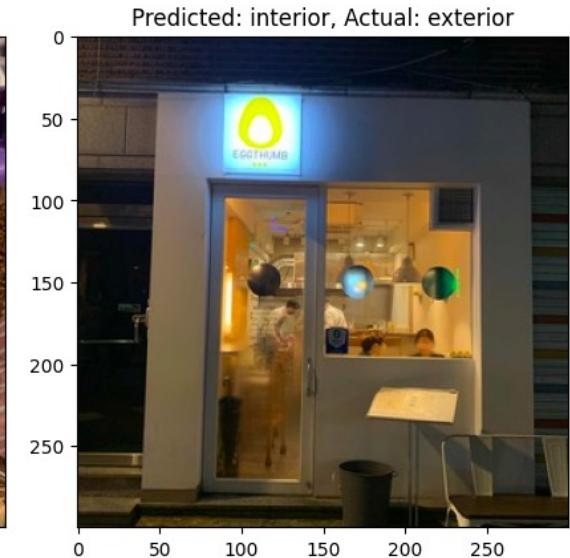
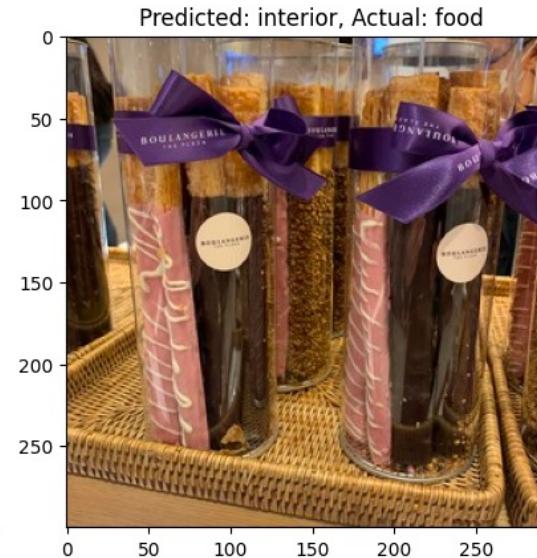
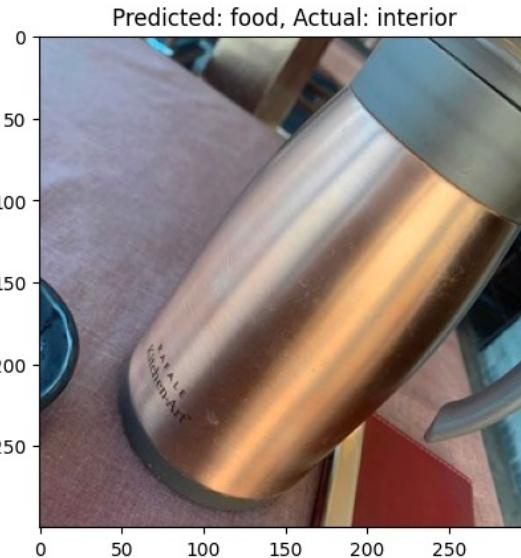
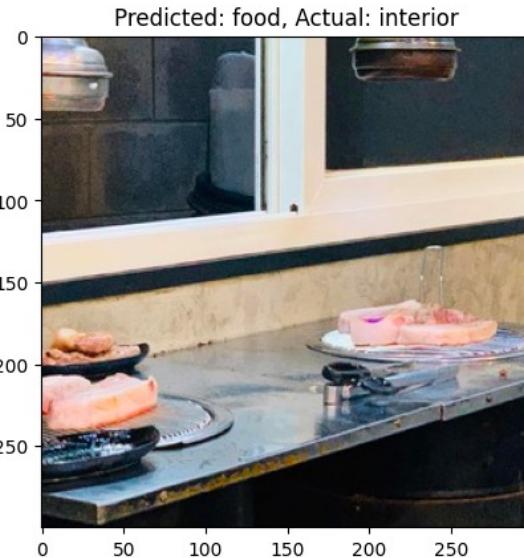
----Correct Predictions---- 10개



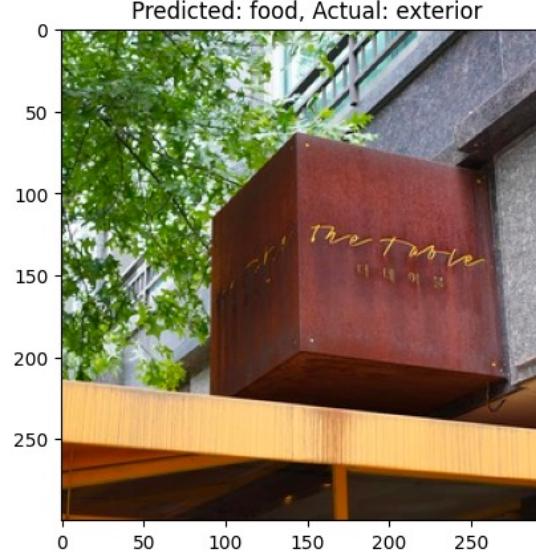
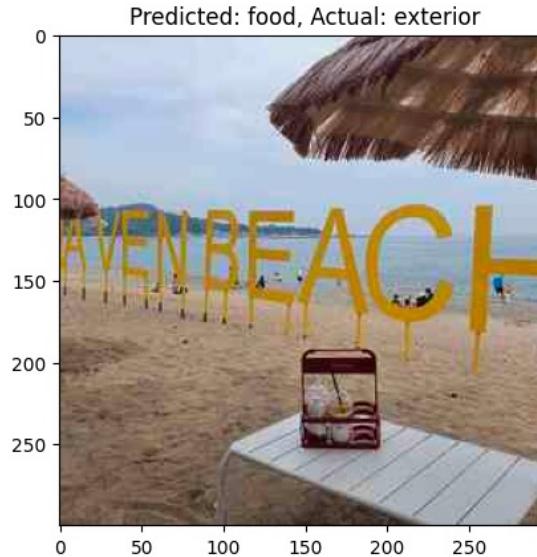
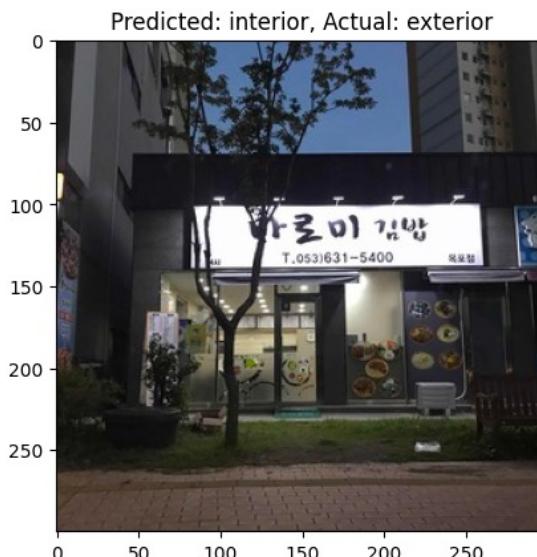
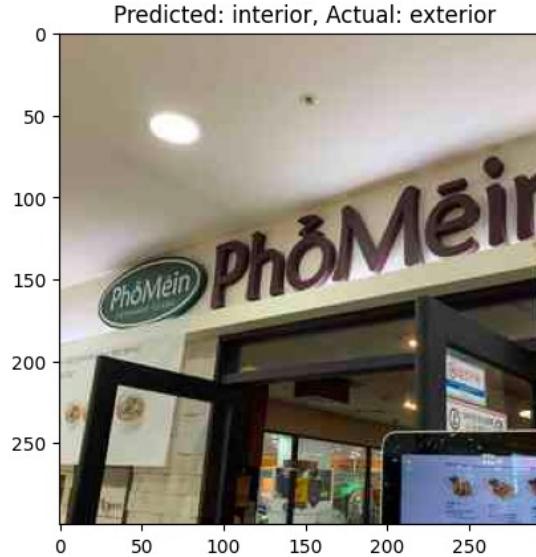
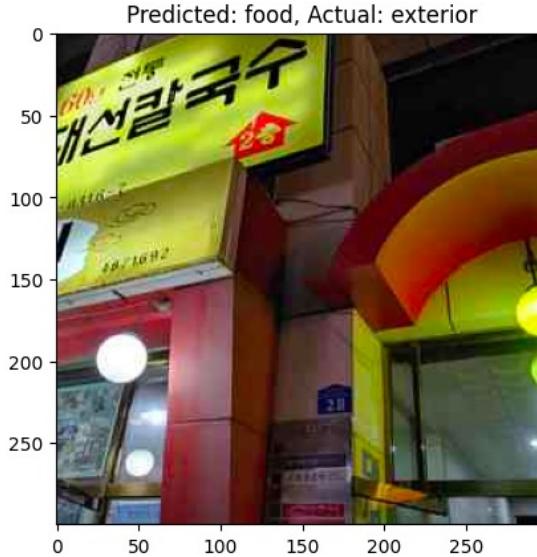
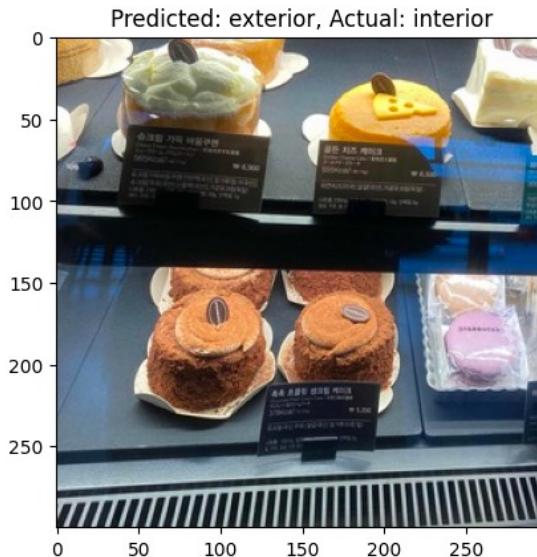
----Correct Predictions----



----Incorrect Predictions---- 10개



----Incorrect Predictions----



Accuracy & Confusion Matrix

Accuracy : 0.8107, FC_Only 때보다는 준수한 정확도, 여러 번의 cnn node 값 수정 trials의 결과

Accuracy = 0.8107

Test Loss: 0.4804

Confusion Matrix:

[[195 35 93]

[3 600 29]

[32 80 370]]

Building Block Method

ResNet50(weights=ResNet50_Weights.DEFAULT)

Replace each 2-layer block in the 34-layer net with 3-layer bottleneck block, resulting in a 50-layer ResNet

```
ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  (layer1): Sequential(
    (0): Bottleneck(
      (conv1): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (downsample): Sequential(
        (0): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
    (1): Bottleneck(
      (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    ...
  )
  (avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
  (fc): Linear(in_features=2048, out_features=1000, bias=True)
)
```

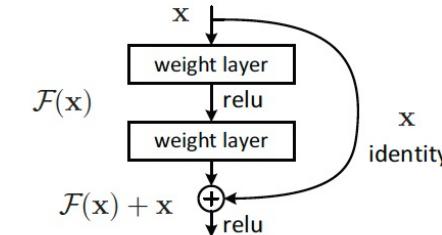


Figure 2. Residual learning: a building block.

layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
		7x7, 64, stride 2	3x3 max pool, stride 2	1x1, 64 3x3, 64 1x1, 256	1x1, 64 3x3, 64 1x1, 256	1x1, 64 3x3, 64 1x1, 256
conv1	112x112					
conv2_x	56x56	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
conv3_x	28x28	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 8$
conv4_x	14x14	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$
conv5_x	7x7	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
	1x1				average pool, 1000-d fc, softmax	
FLOPs		1.8×10^9	3.6×10^9	3.8×10^9	7.6×10^9	11.3×10^9

nn.Module

Pre-Trained model의 features를 identity matrix로 받고, FC layer로 연결, fc1은 requires_grad -> FALSE

이후 fc2, fc3는 모두 TRUE로 받고 training을 진행, 여러 번의 시행 결과 dropout을 제외했을 때 좋은 성능을 보임.

```
class building_block_method(nn.Module):
    def __init__(self, num_classes, fc_requires_grad=True):
        super(building_block_method, self).__init__()

        resnet_model = resnet50(weights=ResNet50_Weights.DEFAULT)
        in_features = resnet_model.fc.in_features
        resnet_model.fc = nn.Identity()

        for param in resnet_model.parameters():
            param.requires_grad = False

        self.resnet = resnet_model
        self.fc1 = nn.Linear(in_features, 32)
        self.fc1.requires_grad = False
        self.fc2 = nn.Linear(32, 64)
        self.fc2.requires_grad = fc_requires_grad
        self.fc3 = nn.Linear(64, num_classes)
        self.fc3.requires_grad = fc_requires_grad

    def forward(self, x):
        x = self.resnet(x)
        x = torch.relu(self.fc1(x))
        x = torch.relu(self.fc2(x))
        x = self.fc3(x)

        return x
```

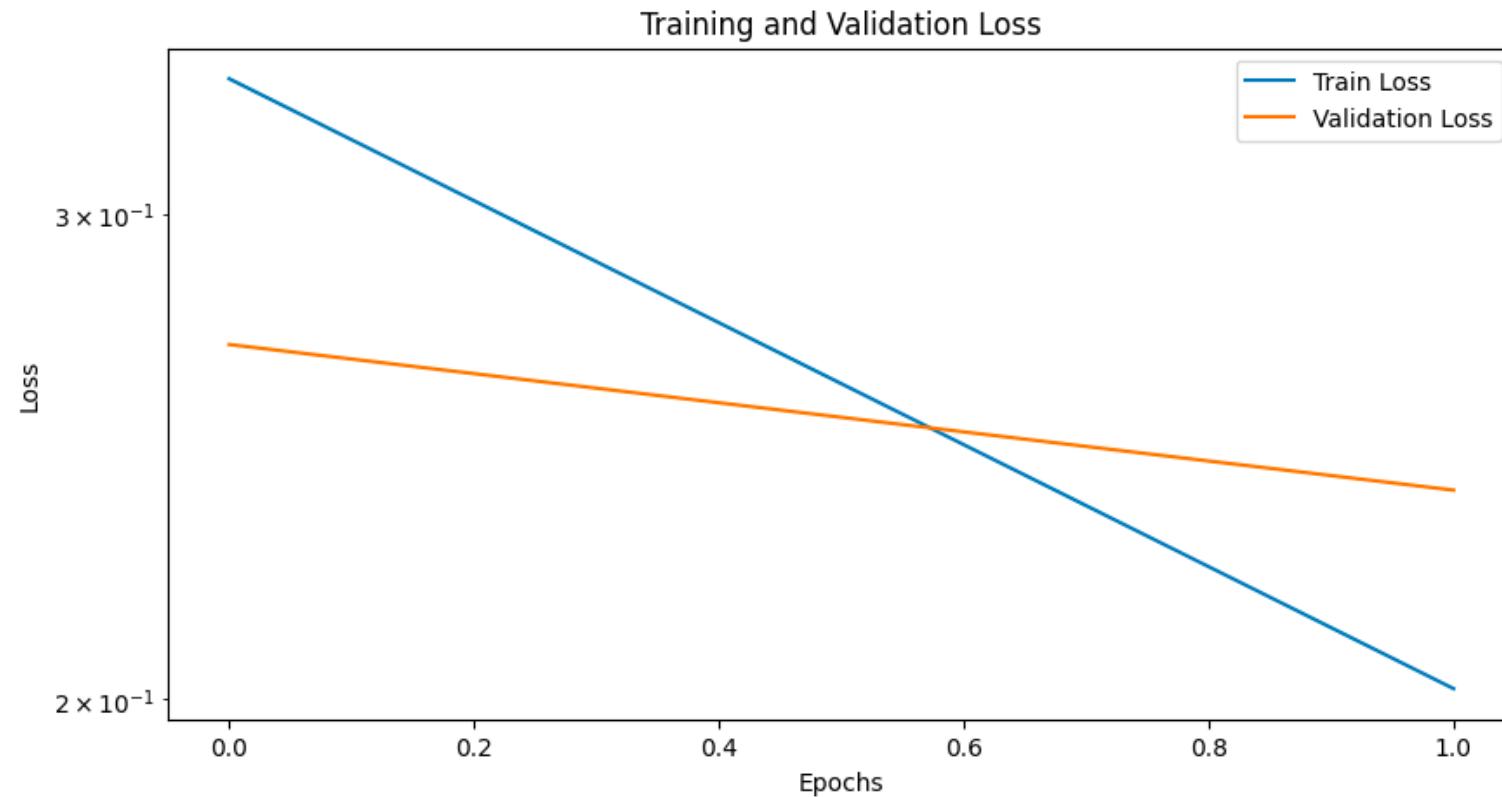
Pre-Trained

Plot Loss

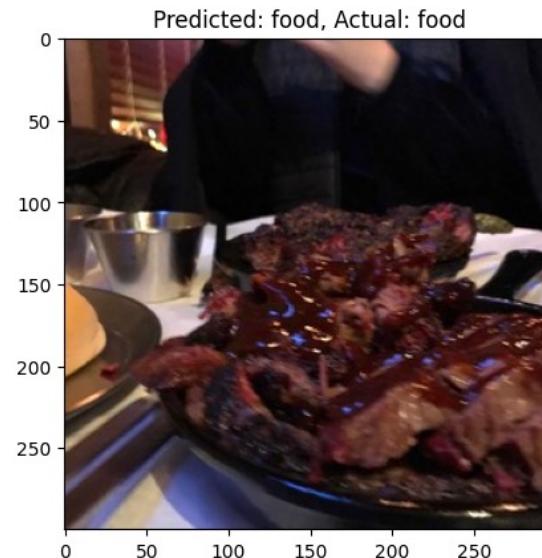
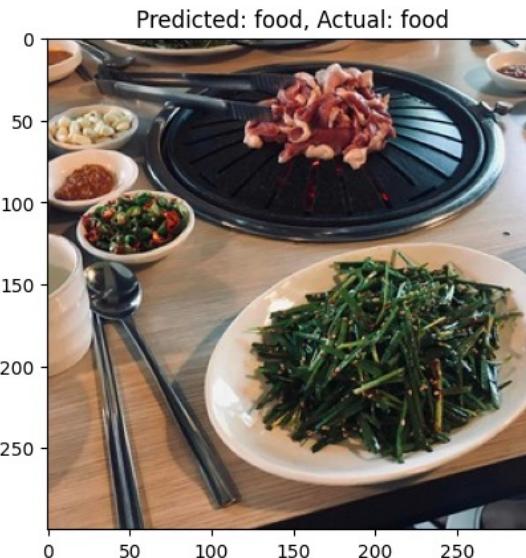
전체적으로 준수하게 감소된 Loss 수준을 보임.

Epoch 1/2 – Train Loss: 0.3357, Validation Loss: 0.2690

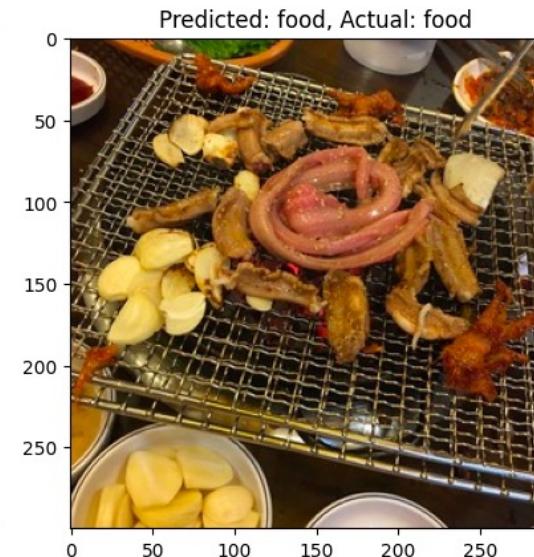
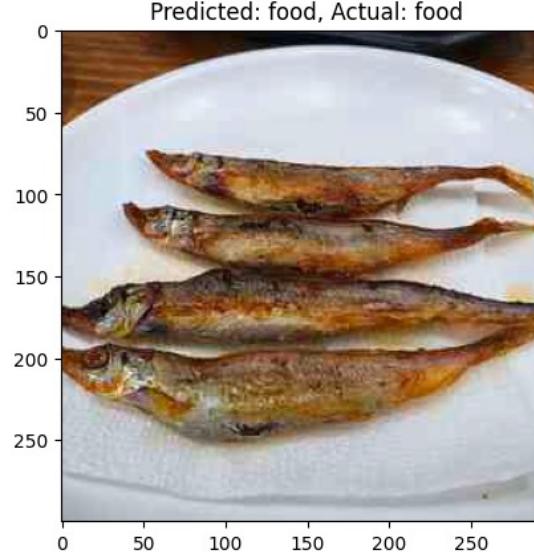
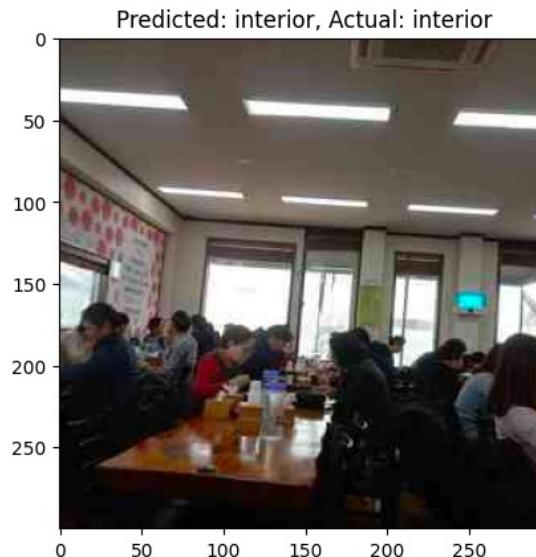
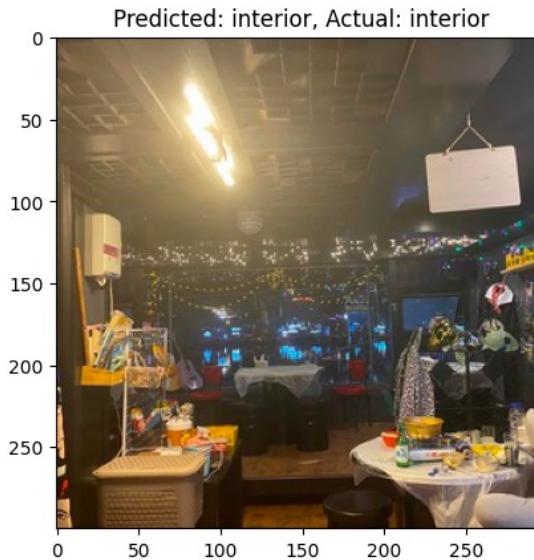
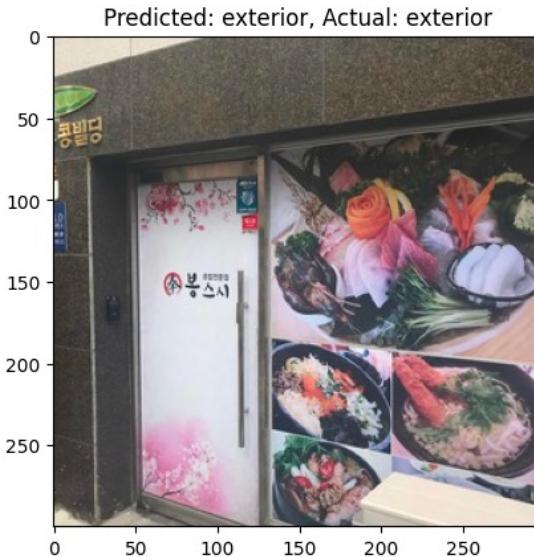
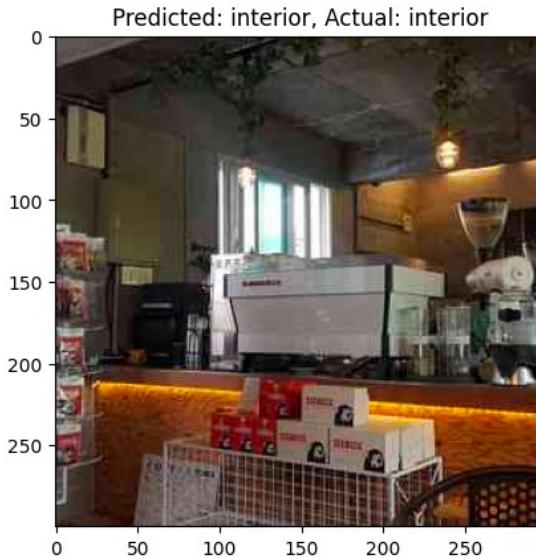
Epoch 2/2 – Train Loss: 0.2018, Validation Loss: 0.2382



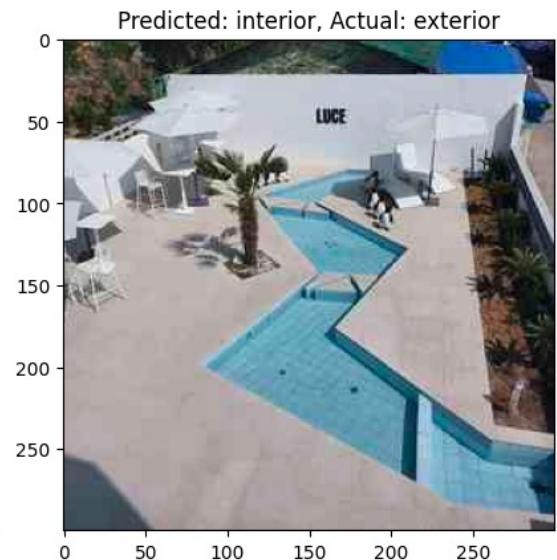
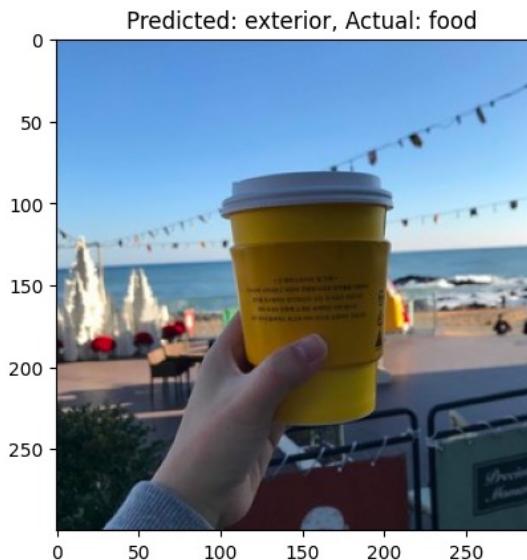
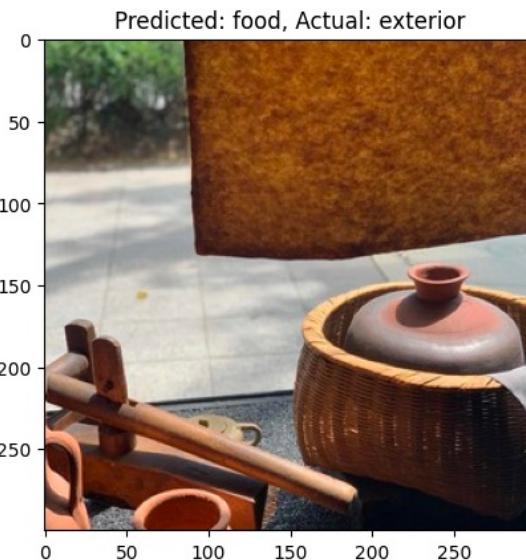
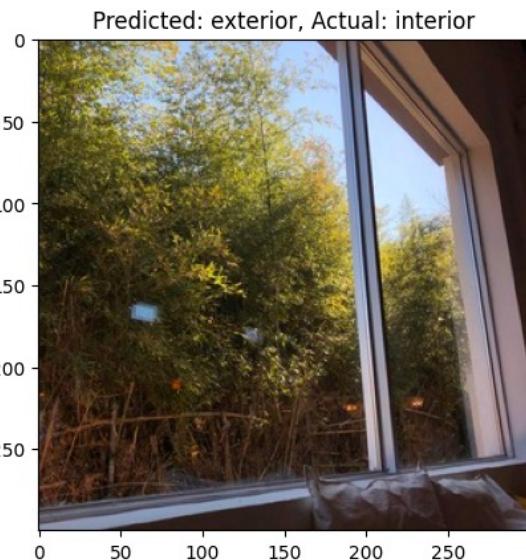
----Correct Predictions---- 10개



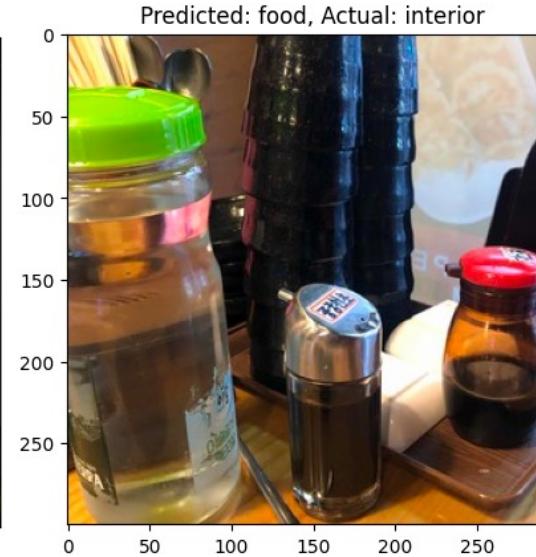
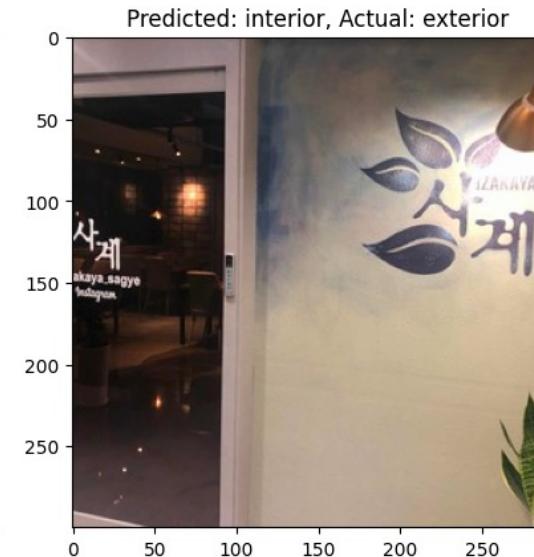
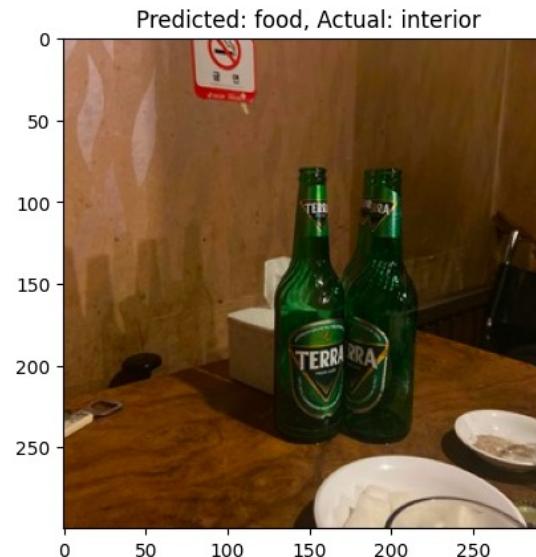
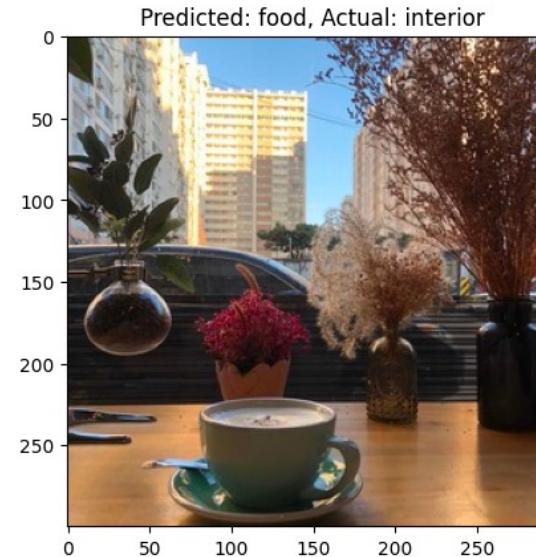
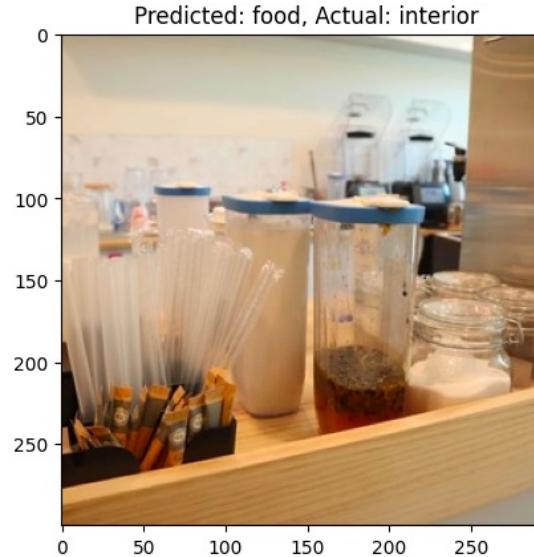
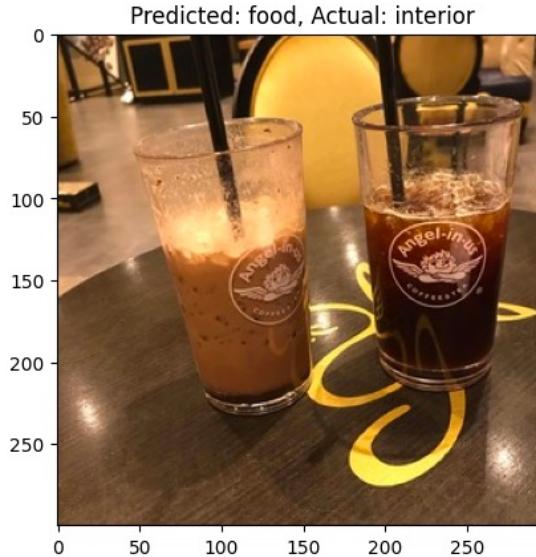
----Correct Predictions----



----Incorrect Predictions---- 10개



----Incorrect Predictions----



Accuracy & Confusion Matrix

Accuracy : 0.9374, 세 개의 method 중 가장 준수한 정확도,

추가적으로, ResNet50_Weights.DEFAULT을 활용한 building block method는
resnet50 구조의 전체 parameter를 학습 했을 때의 accuracy와 거의 유사한 값을 산출했음.
따라서, Generalization이 매우 잘 되어있는 Parameter임을 확인할 수 있음.

```
Accuracy = 0.9374
Test Loss: 0.1725
Confusion Matrix:
[[269    2   33]
 [ 1 629    6]
 [ 18   30 449]]
```

