

# Results

December 19, 2024

```
[ ]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix

import torch
import torchvision.models
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data.dataloader import DataLoader
from torchvision.datasets import ImageFolder
from torch.utils.data import random_split
import torchvision.transforms as transforms
from tqdm import tqdm
import torch.optim as optim

from torchvision.models import resnet50, ResNet50_Weights
import ssl
ssl._create_default_https_context = ssl._create_unverified_context
```

```
[ ]: dir(torchvision.models)
```

```
[ ]: ['AlexNet',
      'AlexNet_Weights',
      'ConvNeXt',
      'ConvNeXt_Base_Weights',
      'ConvNeXt_Large_Weights',
      'ConvNeXt_Small_Weights',
      'ConvNeXt_Tiny_Weights',
      'DenseNet',
      'DenseNet121_Weights',
      'DenseNet161_Weights',
      'DenseNet169_Weights',
      'DenseNet201_Weights',
      'EfficientNet',
      'EfficientNet_B0_Weights',
      'EfficientNet_B1_Weights',
      'EfficientNet_B2_Weights',
```

'EfficientNet\_B3\_Weights',  
'EfficientNet\_B4\_Weights',  
'EfficientNet\_B5\_Weights',  
'EfficientNet\_B6\_Weights',  
'EfficientNet\_B7\_Weights',  
'EfficientNet\_V2\_L\_Weights',  
'EfficientNet\_V2\_M\_Weights',  
'EfficientNet\_V2\_S\_Weights',  
'GoogLeNet',  
'GoogLeNetOutputs',  
'GoogLeNet\_Weights',  
'Inception3',  
'InceptionOutputs',  
'Inception\_V3\_Weights',  
'MNASNet',  
'MNASNet0\_5\_Weights',  
'MNASNet0\_75\_Weights',  
'MNASNet1\_0\_Weights',  
'MNASNet1\_3\_Weights',  
'MaxVit',  
'MaxVit\_T\_Weights',  
'MobileNetV2',  
'MobileNetV3',  
'MobileNet\_V2\_Weights',  
'MobileNet\_V3\_Large\_Weights',  
'MobileNet\_V3\_Small\_Weights',  
'RegNet',  
'RegNet\_X\_16GF\_Weights',  
'RegNet\_X\_1\_6GF\_Weights',  
'RegNet\_X\_32GF\_Weights',  
'RegNet\_X\_3\_2GF\_Weights',  
'RegNet\_X\_400MF\_Weights',  
'RegNet\_X\_800MF\_Weights',  
'RegNet\_X\_8GF\_Weights',  
'RegNet\_Y\_128GF\_Weights',  
'RegNet\_Y\_16GF\_Weights',  
'RegNet\_Y\_1\_6GF\_Weights',  
'RegNet\_Y\_32GF\_Weights',  
'RegNet\_Y\_3\_2GF\_Weights',  
'RegNet\_Y\_400MF\_Weights',  
'RegNet\_Y\_800MF\_Weights',  
'RegNet\_Y\_8GF\_Weights',  
'ResNeXt101\_32X8D\_Weights',  
'ResNeXt101\_64X4D\_Weights',  
'ResNeXt50\_32X4D\_Weights',  
'ResNet',  
'ResNet101\_Weights',

'ResNet152\_Weights',  
'ResNet18\_Weights',  
'ResNet34\_Weights',  
'ResNet50\_Weights',  
'ShuffleNetV2',  
'ShuffleNet\_V2\_X0\_5\_Weights',  
'ShuffleNet\_V2\_X1\_0\_Weights',  
'ShuffleNet\_V2\_X1\_5\_Weights',  
'ShuffleNet\_V2\_X2\_0\_Weights',  
'SqueezeNet',  
'SqueezeNet1\_0\_Weights',  
'SqueezeNet1\_1\_Weights',  
'SwinTransformer',  
'Swin\_B\_Weights',  
'Swin\_S\_Weights',  
'Swin\_T\_Weights',  
'Swin\_V2\_B\_Weights',  
'Swin\_V2\_S\_Weights',  
'Swin\_V2\_T\_Weights',  
'VGG',  
'VGG11\_BN\_Weights',  
'VGG11\_Weights',  
'VGG13\_BN\_Weights',  
'VGG13\_Weights',  
'VGG16\_BN\_Weights',  
'VGG16\_Weights',  
'VGG19\_BN\_Weights',  
'VGG19\_Weights',  
'ViT\_B\_16\_Weights',  
'ViT\_B\_32\_Weights',  
'ViT\_H\_14\_Weights',  
'ViT\_L\_16\_Weights',  
'ViT\_L\_32\_Weights',  
'VisionTransformer',  
'Weights',  
'WeightsEnum',  
'Wide\_ResNet101\_2\_Weights',  
'Wide\_ResNet50\_2\_Weights',  
'\_GoogLeNetOutputs',  
'\_InceptionOutputs',  
'\_\_builtins\_\_',  
'\_\_cached\_\_',  
'\_\_doc\_\_',  
'\_\_file\_\_',  
'\_\_loader\_\_',  
'\_\_name\_\_',  
'\_\_package\_\_',

```
'__path__',  
'__spec__',  
'_api',  
'_meta',  
'_utils',  
'alexnet',  
'convnext',  
'convnext_base',  
'convnext_large',  
'convnext_small',  
'convnext_tiny',  
'densenet',  
'densenet121',  
'densenet161',  
'densenet169',  
'densenet201',  
'detection',  
'efficientnet',  
'efficientnet_b0',  
'efficientnet_b1',  
'efficientnet_b2',  
'efficientnet_b3',  
'efficientnet_b4',  
'efficientnet_b5',  
'efficientnet_b6',  
'efficientnet_b7',  
'efficientnet_v2_l',  
'efficientnet_v2_m',  
'efficientnet_v2_s',  
'get_model',  
'get_model_builder',  
'get_model_weights',  
'get_weight',  
'googlenet',  
'inception',  
'inception_v3',  
'list_models',  
'maxvit',  
'maxvit_t',  
'mnasnet',  
'mnasnet0_5',  
'mnasnet0_75',  
'mnasnet1_0',  
'mnasnet1_3',  
'mobilenet',  
'mobilenet_v2',  
'mobilenet_v3_large',
```

'mobilenet\_v3\_small',  
'mobilenetv2',  
'mobilenetv3',  
'optical\_flow',  
'quantization',  
'regnet',  
'regnet\_x\_16gf',  
'regnet\_x\_1\_6gf',  
'regnet\_x\_32gf',  
'regnet\_x\_3\_2gf',  
'regnet\_x\_400mf',  
'regnet\_x\_800mf',  
'regnet\_x\_8gf',  
'regnet\_y\_128gf',  
'regnet\_y\_16gf',  
'regnet\_y\_1\_6gf',  
'regnet\_y\_32gf',  
'regnet\_y\_3\_2gf',  
'regnet\_y\_400mf',  
'regnet\_y\_800mf',  
'regnet\_y\_8gf',  
'resnet',  
'resnet101',  
'resnet152',  
'resnet18',  
'resnet34',  
'resnet50',  
'resnext101\_32x8d',  
'resnext101\_64x4d',  
'resnext50\_32x4d',  
'segmentation',  
'shufflenet\_v2\_x0\_5',  
'shufflenet\_v2\_x1\_0',  
'shufflenet\_v2\_x1\_5',  
'shufflenet\_v2\_x2\_0',  
'shufflenetv2',  
'squeezenet',  
'squeezenet1\_0',  
'squeezenet1\_1',  
'swin\_b',  
'swin\_s',  
'swin\_t',  
'swin\_transformer',  
'swin\_v2\_b',  
'swin\_v2\_s',  
'swin\_v2\_t',  
'vgg',

```

'vgg11',
'vgg11_bn',
'vgg13',
'vgg13_bn',
'vgg16',
'vgg16_bn',
'vgg19',
'vgg19_bn',
'video',
'vision_transformer',
'vit_b_16',
'vit_b_32',
'vit_h_14',
'vit_l_16',
'vit_l_32',
'wide_resnet101_2',
'wide_resnet50_2']

```

```
[ ]: resnet50(weights=ResNet50_Weights.DEFAULT)
```

```

[ ]: ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (relu): ReLU(inplace=True)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1,
ceil_mode=False)
  (layer1): Sequential(
    (0): Bottleneck(
      (conv1): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (downsample): Sequential(
        (0): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
  )
)

```

```

        (1): Bottleneck(
          (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
          (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
          (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
          (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (relu): ReLU(inplace=True)
        )
      (2): Bottleneck(
        (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
      )
    )
  (layer2): Sequential(
    (0): Bottleneck(
      (conv1): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (downsample): Sequential(
        (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
  )
  (1): Bottleneck(

```

```

        (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
    )
    (2): Bottleneck(
        (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
    )
    (3): Bottleneck(
        (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
    )
    )
    (layer3): Sequential(
        (0): Bottleneck(
            (conv1): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
            (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,

```



```

track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (downsample): Sequential(
      (0): Conv2d(512, 1024, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (1): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
  )
  (2): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
  )
  (3): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,

```

```

track_running_stats=True)
    (relu): ReLU(inplace=True)
)
(4): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (relu): ReLU(inplace=True)
)
(5): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (relu): ReLU(inplace=True)
)
)
(layer4): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (downsample): Sequential(
      (0): Conv2d(1024, 2048, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

```

    )
    )
    (1): Bottleneck(
      (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
    (2): Bottleneck(
      (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
  )
  (avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
  (fc): Linear(in_features=2048, out_features=1000, bias=True)
)

```

```

[ ]: class ImageClassificationModel_FC(nn.Module):
    def __init__(self):
        super(ImageClassificationModel_FC, self).__init__()
        self.fc1 = nn.Linear(3 * 300 * 300, 256)
        self.fc2 = nn.Linear(256, 128)
        self.fc3 = nn.Linear(128, 128)
        self.fc4 = nn.Linear(128, 128)
        self.fc5 = nn.Linear(128, 128)
        self.fc6 = nn.Linear(128, 128)

        self.dropout = nn.Dropout(0.5)

        self.fcm = nn.Linear(128, 3)

```

```

def forward(self, x):
    x = x.reshape(x.size(0), -1)
    x = torch.relu(self.fc1(x))
    x = torch.relu(self.fc2(x))
    x = torch.relu(self.fc3(x))
    x = torch.relu(self.fc4(x))
    x = torch.relu(self.fc5(x))
    x = torch.relu(self.fc6(x))

    x = self.dropout(x)

    x = self.fcm(x)

    return x

```

```

[ ]: class ImageClassificationModel_CNN(nn.Module):
    def __init__(self):
        super(ImageClassificationModel_CNN, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, kernel_size=3, stride=1, padding=1)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1)
        self.conv3 = nn.Conv2d(64, 64, kernel_size=3, stride=1, padding=1)

        self.pool = nn.MaxPool2d(kernel_size=2, stride=2, padding = 0)
        self.dropout = nn.Dropout(0.5)

        self.fc1 = nn.Linear(64 * 37 * 37, 64)
        self.fc2 = nn.Linear(64, 32)
        self.fc3 = nn.Linear(32, 3)

    def forward(self, x):
        x = self.pool(torch.relu(self.conv1(x)))
        x = self.pool(torch.relu(self.conv2(x)))
        x = self.pool(torch.relu(self.conv3(x)))

        x = x.view(-1, 64 * 37 * 37)

        x = self.dropout(x)

        x = torch.relu(self.fc1(x))
        x = torch.relu(self.fc2(x))
        x = self.fc3(x)

        return x

```

```

[ ]: class building_block_method(nn.Module):
    def __init__(self, num_classes, fc_requires_grad=True):
        super(building_block_method, self).__init__()

```

```

resnet_model = resnet50(weights=ResNet50_Weights.DEFAULT)
in_features = resnet_model.fc.in_features
resnet_model.fc = nn.Identity()

for param in resnet_model.parameters():
    param.requires_grad = False

self.resnet = resnet_model
self.fc1 = nn.Linear(in_features, 32)
self.fc1.requires_grad = False
self.fc2 = nn.Linear(32, 64)
self.fc2.requires_grad = fc_requires_grad
self.fc3 = nn.Linear(64, num_classes)
self.fc3.requires_grad = fc_requires_grad

def forward(self,x):
    x = self.resnet(x)
    x = torch.relu(self.fc1(x))
    x = torch.relu(self.fc2(x))
    x = self.fc3(x)

    return x

```

```

[ ]: class ImageDataLoader:
    def __init__(self):
        transform = transforms.Compose([
            transforms.Resize((300, 300)),
            transforms.ToTensor(),
            transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.
↪225])
        ])

        root_dir = 'ProjectImages'
        full_train_dataset = ImageFolder(root=root_dir, transform=transform)
        self.classes = full_train_dataset.classes

        total_size = len(full_train_dataset)
        train_size = int(total_size * 0.7)
        val_size = int(total_size * 0.1)
        test_size = total_size - train_size - val_size

        self.train_data, self.val_data, self.test_data =
↪random_split(full_train_dataset, [train_size, val_size, test_size])

        self.train_loader = DataLoader(self.train_data, batch_size=32,
↪shuffle=True, num_workers=2, pin_memory=True)

```

```

        self.val_loader = DataLoader(self.val_data, batch_size=32,
↪shuffle=True, num_workers=2, pin_memory=True)
        self.test_loader = DataLoader(self.test_data, batch_size=32,
↪shuffle=True, num_workers=2, pin_memory=True)

```

```

[ ]: class ModelTrainer:
    def __init__(self, model, train_loader, val_loader, test_loader, classes,
↪epochs=10, learning_rate=0.001):
        self.model = model
        self.train_loader = train_loader
        self.val_loader = val_loader
        self.test_loader = test_loader
        self.classes = classes
        self.criterion = nn.CrossEntropyLoss()
        self.optimizer = optim.Adam(model.parameters(), lr=learning_rate)

        self.epochs = epochs

        self.train_losses = []
        self.val_losses = []

    def train(self, device):
        self.model.to(device)
        for epoch in range(self.epochs):
            self.model.train()
            train_loss = 0
            with tqdm(total=len(self.train_loader), desc=f"Epoch {epoch + 1}/
↪{self.epochs}", unit="batch", leave=False) as pbar:
                for features, labels in self.train_loader:
                    features, labels = features.to(device), labels.to(device)
                    outputs = self.model(features)
                    loss = self.criterion(outputs, labels)
                    train_loss += loss.item()

                    self.optimizer.zero_grad()
                    loss.backward()
                    self.optimizer.step()

                pbar.set_postfix({'train_loss': f'{loss.item():.4f}'})
                pbar.update(1)

            train_loss /= len(self.train_loader)
            val_loss = self.validate(device)

            self.train_losses.append(train_loss)
            self.val_losses.append(val_loss)

```

```
        print(f"Epoch {epoch + 1}/{self.epochs} - Train Loss: {train_loss:.4f}, Validation Loss: {val_loss:.4f}")
```

```
def validate(self, device):
    self.model.eval()
    val_loss = 0
    with torch.no_grad():
        for features, labels in self.val_loader:
            features, labels = features.to(device), labels.to(device)
            outputs = self.model(features)
            loss = self.criterion(outputs, labels)
            val_loss += loss.item()

    val_loss /= len(self.val_loader)
    return val_loss
```

```
def plot_loss(self):
    plt.figure(figsize=(10, 5))
    plt.plot(self.train_losses, label='Train Loss')
    plt.plot(self.val_losses, label='Validation Loss')
    plt.title('Training and Validation Loss')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.yscale('log')
    plt.legend()
    plt.show()
```

```
def evaluate(self, device):
    self.model.eval()
    test_loss = 0
    all_labels = []
    all_preds = []

    total = 0
    correct = 0

    correct_predictions = []
    incorrect_predictions = []

    with torch.no_grad():
        for features, labels in self.test_loader:
            features, labels = features.to(device), labels.to(device)
            outputs = self.model(features.float())
            loss = self.criterion(outputs, labels)
```

```

        test_loss += loss.item()

        output_prob = F.softmax(outputs, dim=1).cpu().numpy()

        _, predicted = torch.max(outputs, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

        all_labels.extend(labels.cpu().numpy())
        all_preds.extend(predicted.cpu().numpy())

        for f, l, p, o, prob in zip(features, labels, predicted,
↪outputs, output_prob):
            if l == p and len(correct_predictions) < 10:
                correct_predictions.append((f, l, p, o, prob))
            elif l != p and len(incorrect_predictions) < 10:
                incorrect_predictions.append((f, l, p, o, prob))
            if len(correct_predictions) >= 10 and
↪len(incorrect_predictions) >= 10:
                break

        print("\n----- Correct Predictions -----")
        for f, l, p, o, prob in correct_predictions:
            f_permute = f.permute(1, 2, 0).cpu().numpy()

            restored_image = (f_permute * [0.229, 0.224, 0.225]) + [0.485, 0.
↪456, 0.406]
            restored_image = np.clip(restored_image, 0, 1)

            plt.imshow(restored_image, cmap='gray')
            plt.title(f'Predicted: {self.classes[p]}, Actual: {self.
↪classes[l]}')
            plt.show()
            print(f"{o}->{prob}: {self.classes[l]} -> {self.classes[p]}")

        print("\n----- Incorrect Predictions -----")
        for f, l, p, o, prob in incorrect_predictions:
            f_permute = f.permute(1, 2, 0).cpu().numpy()

            restored_image = (f_permute * [0.229, 0.224, 0.225]) + [0.485, 0.
↪456, 0.406]
            restored_image = np.clip(restored_image, 0, 1)

            plt.imshow(restored_image, cmap='gray')
            plt.title(f'Predicted: {self.classes[p]}, Actual: {self.
↪classes[l]}')

```



```

plt.show()
print(f"{o}->{prob}: {self.classes[l]} -> {self.classes[p]}")

test_loss /= len(self.test_loader)
print(f"Accuracy = {correct/total: .4f}")

cm = confusion_matrix(all_labels, all_preds)
print(f"Test Loss: {test_loss:.4f}")
print("Confusion Matrix:")
print(cm)

```

```

[ ]: if __name__ == '__main__':
    device = torch.device('mps') if torch.backends.mps.is_available() else "cpu"
    print(device)

    model1 = ImageClassificationModel_FC()
    model1.to(device)

    model2 = ImageClassificationModel_CNN()
    model2.to(device)

    resnet50(weights=ResNet50_Weights.DEFAULT)
    model3 = building_block_method(num_classes=3, fc_requires_grad=True)
    model3.to(device)

    dataloader = ImageDataLoader()

    trainer1 = ModelTrainer(model1, dataloader.train_loader, dataloader.
    ↪ val_loader, dataloader.test_loader, dataloader.classes, epochs=2)
    trainer1.train(device)
    trainer1.validate(device)
    trainer1.plot_loss()
    trainer1.evaluate(device)

    trainer2 = ModelTrainer(model2, dataloader.train_loader, dataloader.
    ↪ val_loader, dataloader.test_loader, dataloader.classes, epochs=2)
    trainer2.train(device)
    trainer2.validate(device)
    trainer2.plot_loss()
    trainer2.evaluate(device)


    trainer3 = ModelTrainer(model3, dataloader.train_loader, dataloader.
    ↪ val_loader, dataloader.test_loader, dataloader.classes, epochs=2)
    trainer3.train(device)
    trainer3.validate(device)
    trainer3.plot_loss()
    trainer3.evaluate(device)

```

mps

Epoch 1/2 - Train Loss: 0.3357, Validation Loss: 0.2690

Epoch 2/2 - Train Loss: 0.2018, Validation Loss: 0.2382

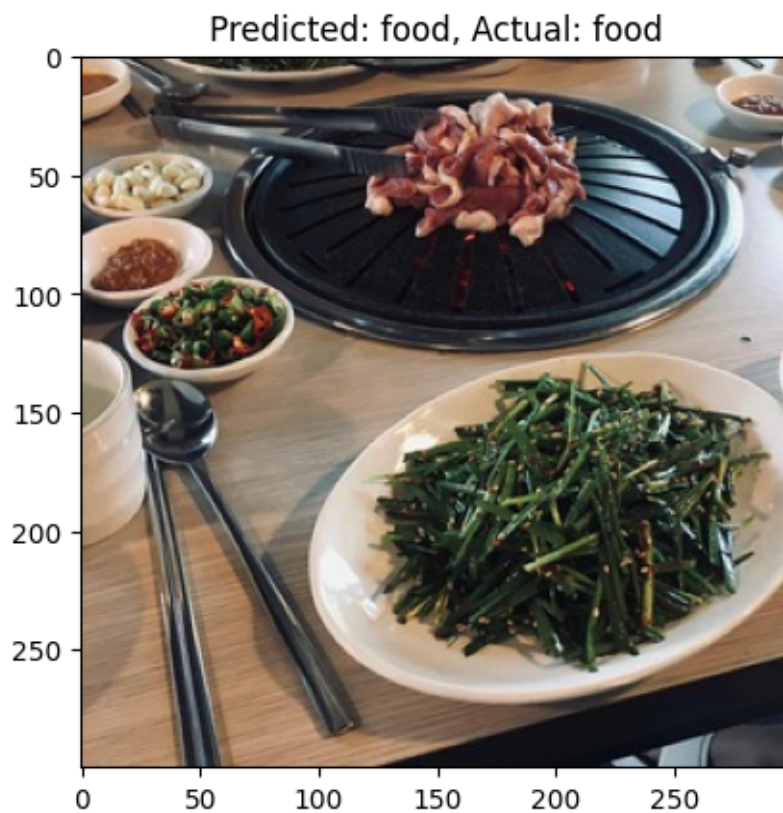


----- Correct Predictions -----

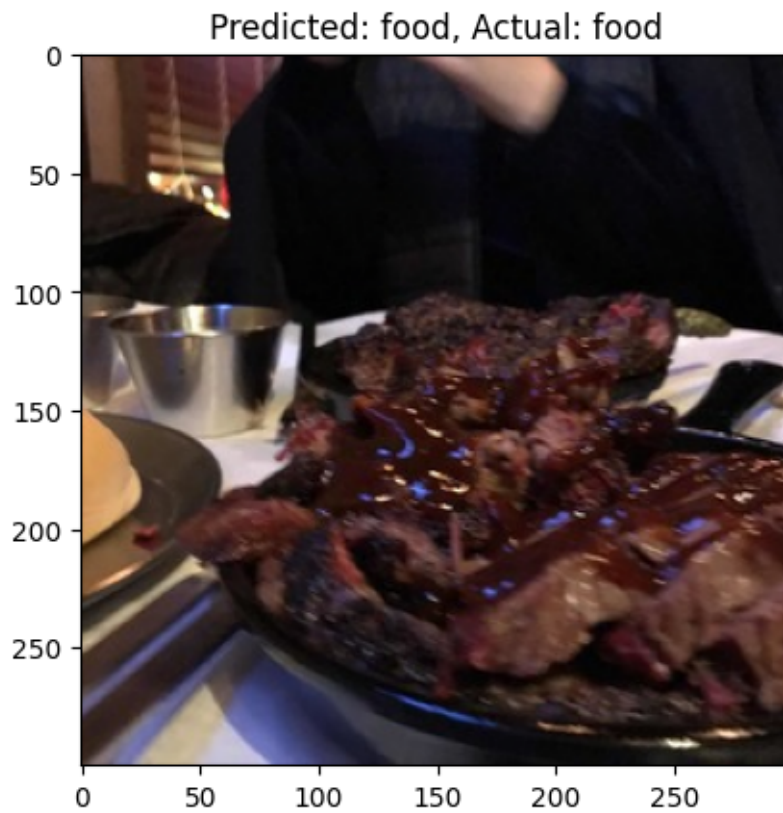
Predicted: exterior, Actual: exterior



```
tensor([ 3.6561, -4.8687, -1.8997], device='mps:0')->[9.9595296e-01  
1.9768013e-04 3.8493443e-03]: exterior -> exterior
```



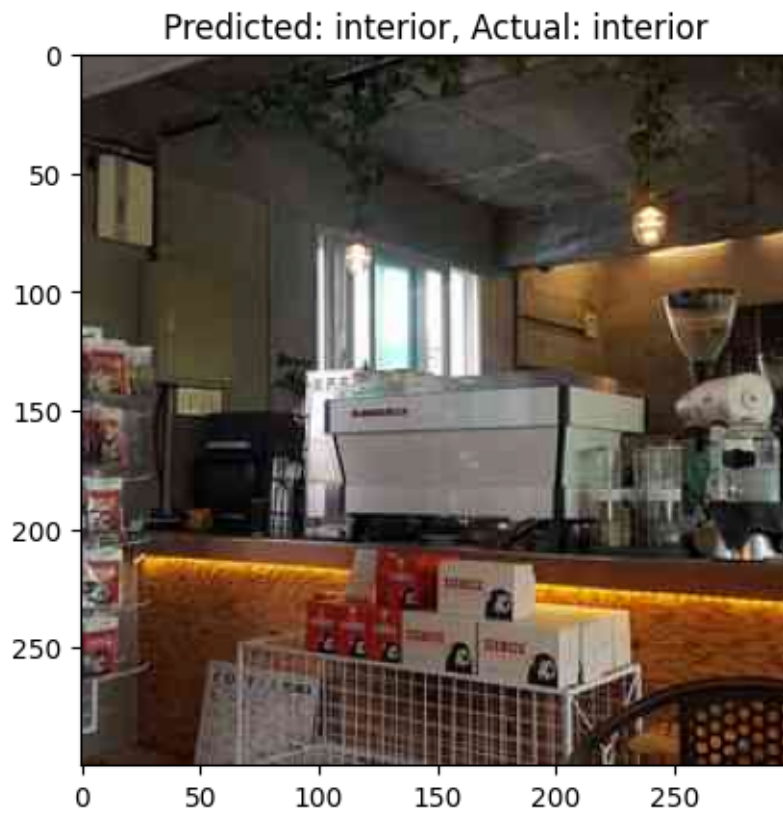
```
tensor([-4.2567,  3.5494, -1.6011], device='mps:0')->[4.0475361e-04  
9.9383461e-01 5.7606036e-03]: food -> food
```



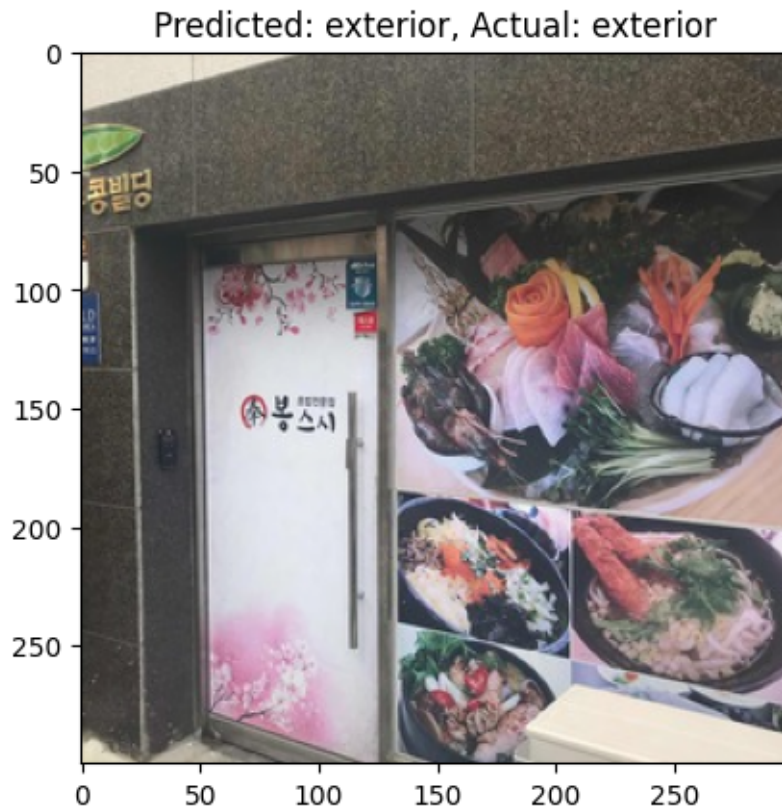
```
tensor([-4.9761,  3.8806, -1.6546], device='mps:0')->[1.4184402e-04  
9.9592870e-01 3.9294548e-03]: food -> food
```



```
tensor([-4.4480,  3.1045, -1.0441], device='mps:0')->[5.1639415e-04  
9.8395073e-01 1.5532908e-02]: food -> food
```




```
tensor([-1.4475, -2.1962,  1.5140], device='mps:0')->[0.04807756 0.02273929  
0.9291832 ]: interior -> interior
```

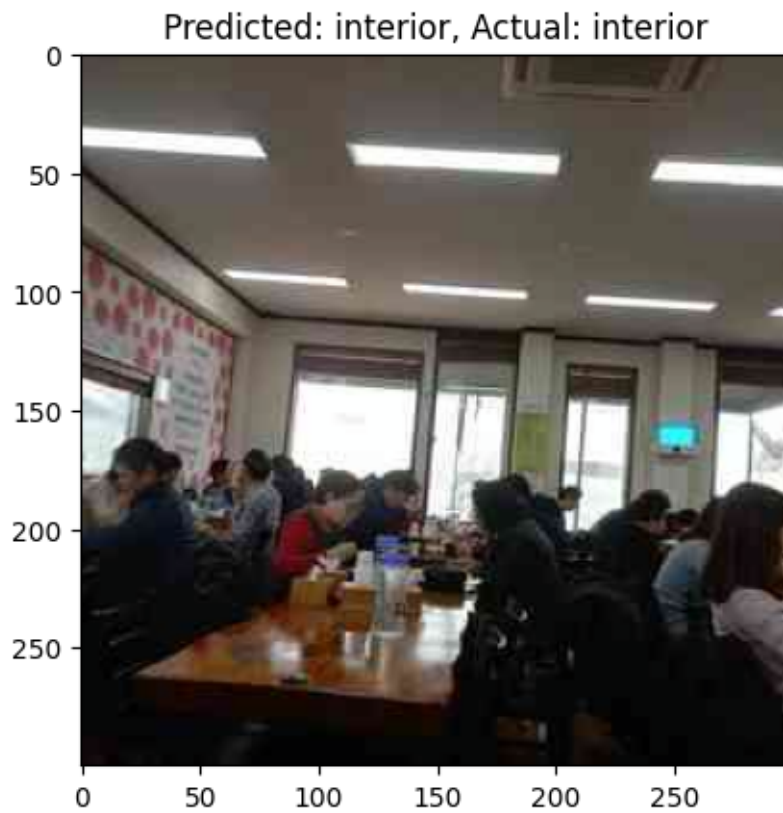


```
tensor([ 0.4371, -2.5603, -0.4303], device='mps:0')->[0.68028396 0.03395645  
0.28575963]: exterior -> exterior
```






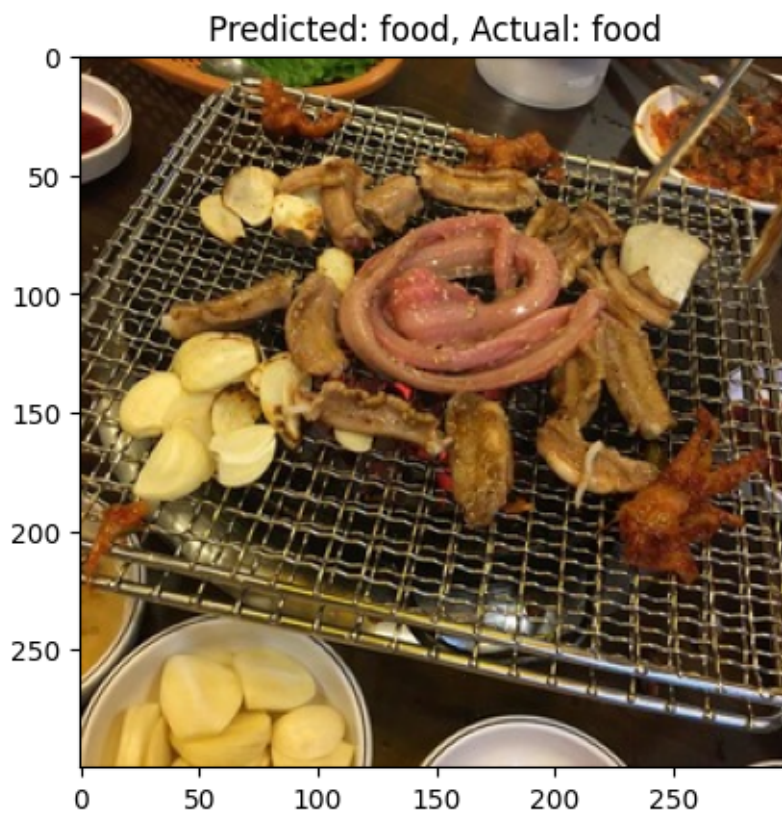
```
tensor([-2.7997, -3.7176,  3.5310], device='mps:0')->[1.7763670e-03  
7.0939941e-04 9.9751425e-01]: interior -> interior
```



```
tensor([-1.9388, -3.2185,  2.7320], device='mps:0')->[0.00925427 0.00257381  
0.9881719 ]: interior -> interior
```

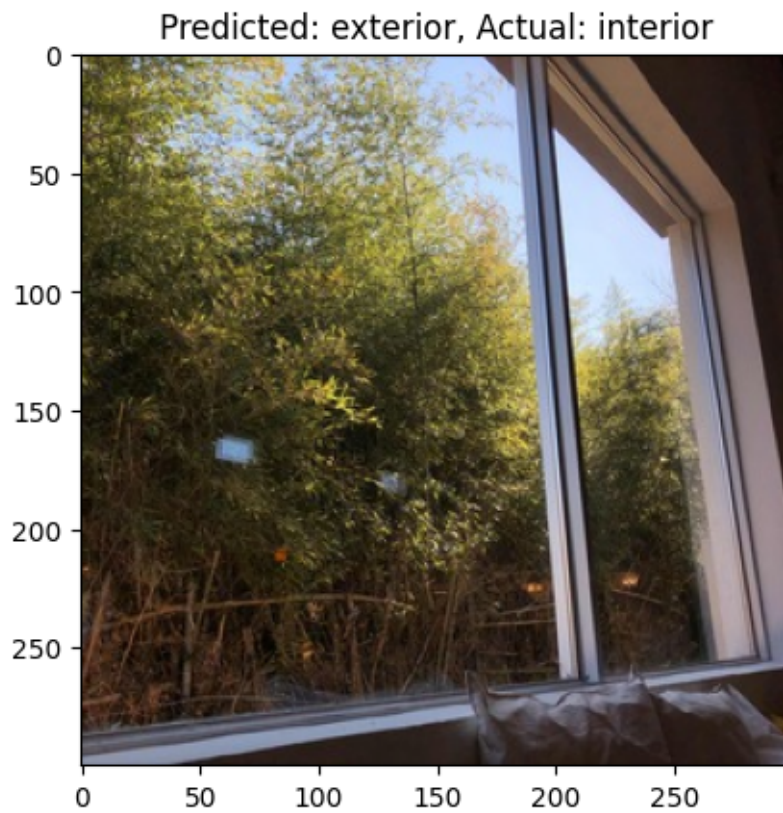


```
tensor([-3.2067,  3.1991, -2.0228], device='mps:0')->[0.00164043 0.99300015  
0.00535938]: food -> food
```

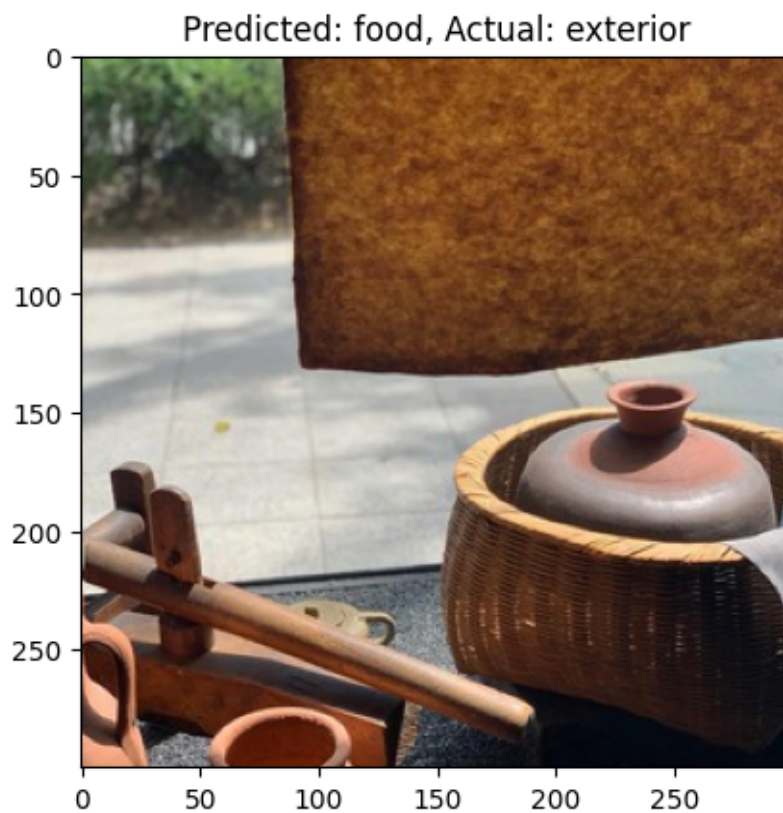


```
tensor([-4.4551,  4.3037, -2.6177], device='mps:0')->[1.5688212e-04  
9.9885786e-01 9.8525907e-04]: food -> food
```

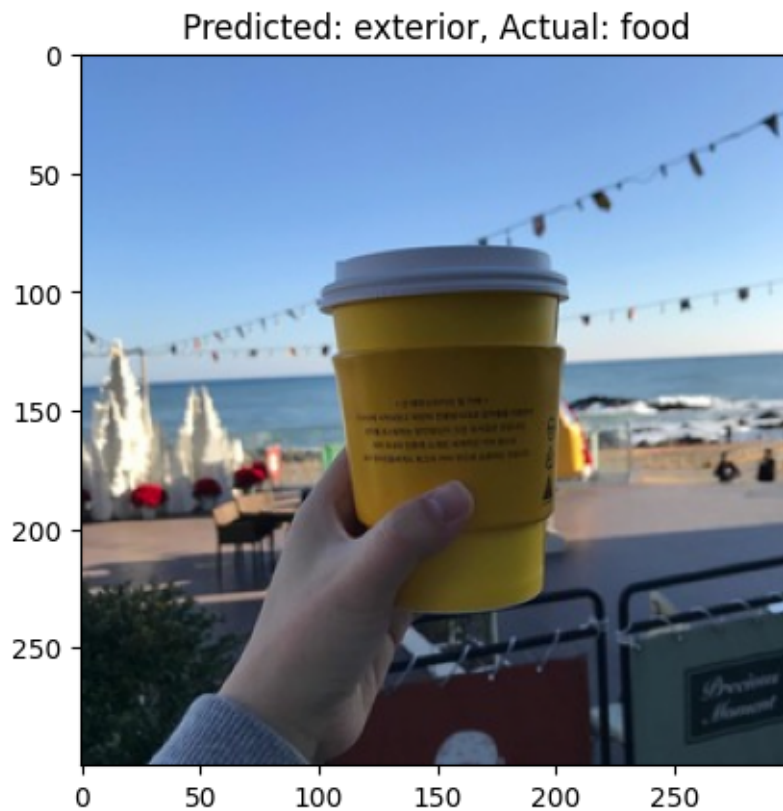
----- Incorrect Predictions -----



```
tensor([ 0.7594, -2.8345,  0.1924], device='mps:0')->[0.6270555  0.01723866  
0.35570583]: interior -> exterior
```




```
tensor([-2.1850,  0.6240, -0.1979], device='mps:0')->[0.04018188 0.66671926  
0.29309884]: exterior -> food
```



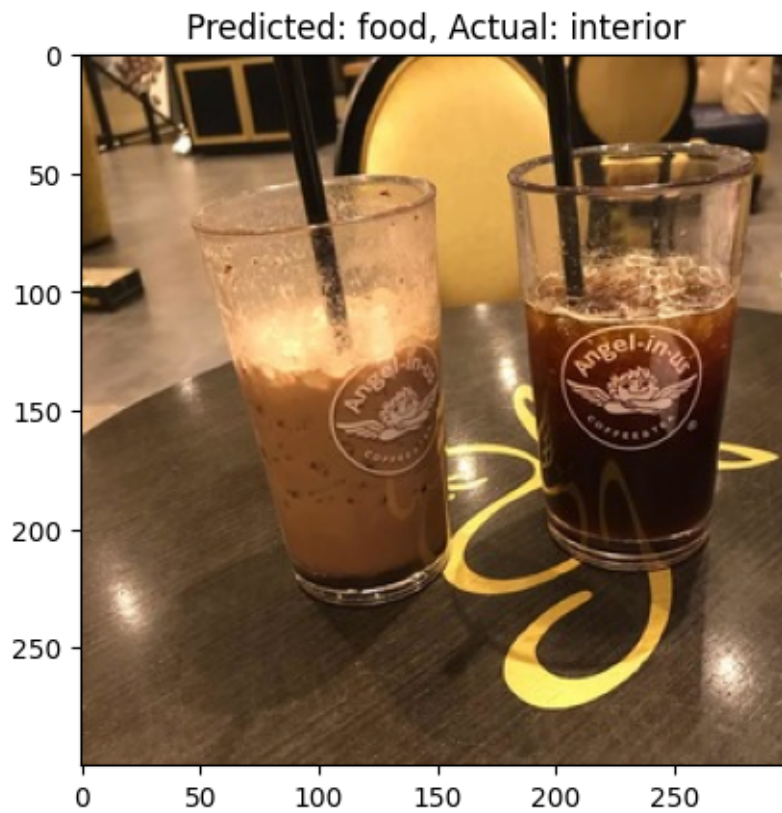
```
tensor([ 0.5220, -1.8772, -0.7506], device='mps:0')->[0.7294539  0.06622441  
0.20432161]: food -> exterior
```



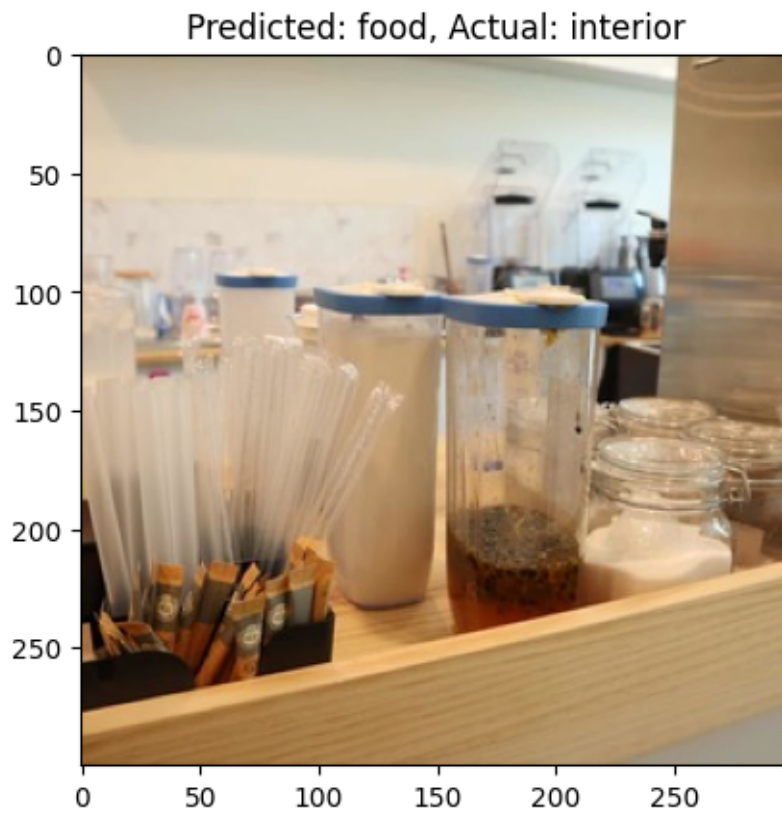


```
tensor([-0.1065, -2.8265,  0.3131], device='mps:0')->[0.38653475 0.02546231  
0.588003  ]: exterior -> interior
```

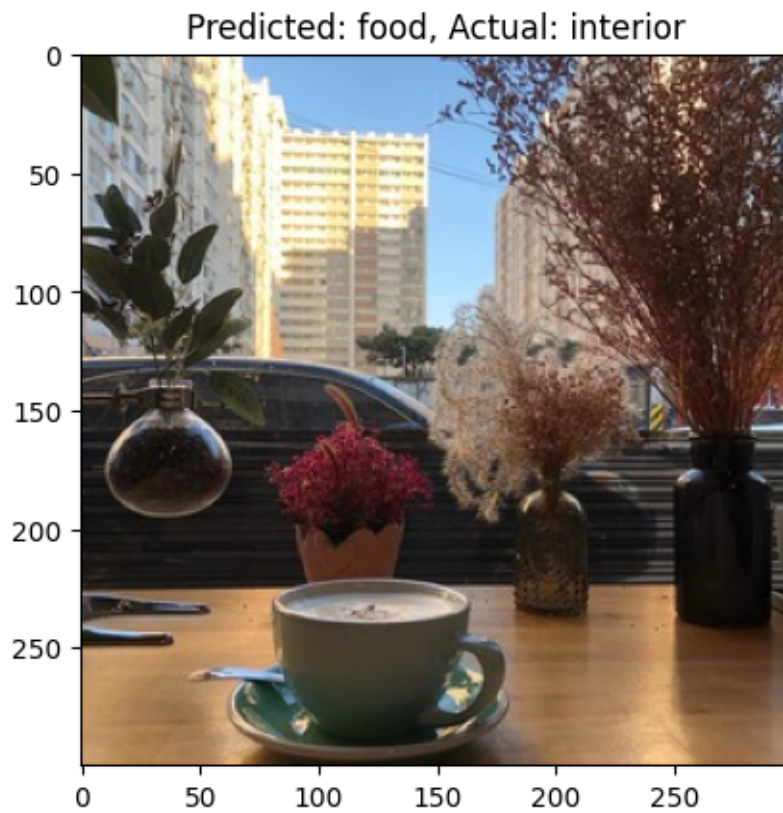





```
tensor([-2.7150,  2.4054, -1.1648], device='mps:0')->[0.00577655 0.9670046  
0.02721885]: interior -> food
```



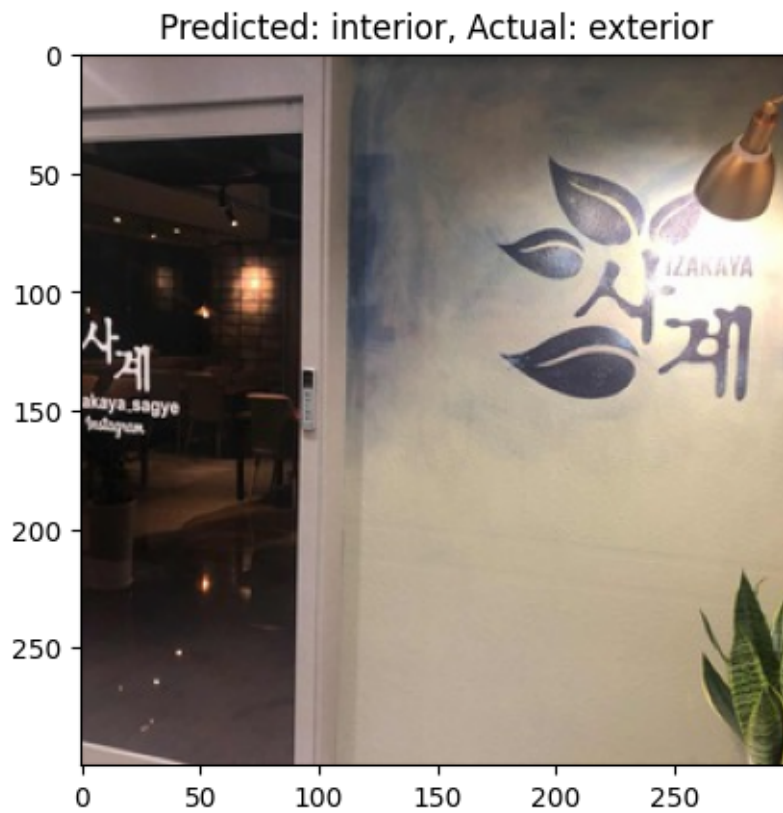
```
tensor([-4.6442,  1.2060,  0.7401], device='mps:0')->[0.00176597 0.6133261  
0.38490793]: interior -> food
```



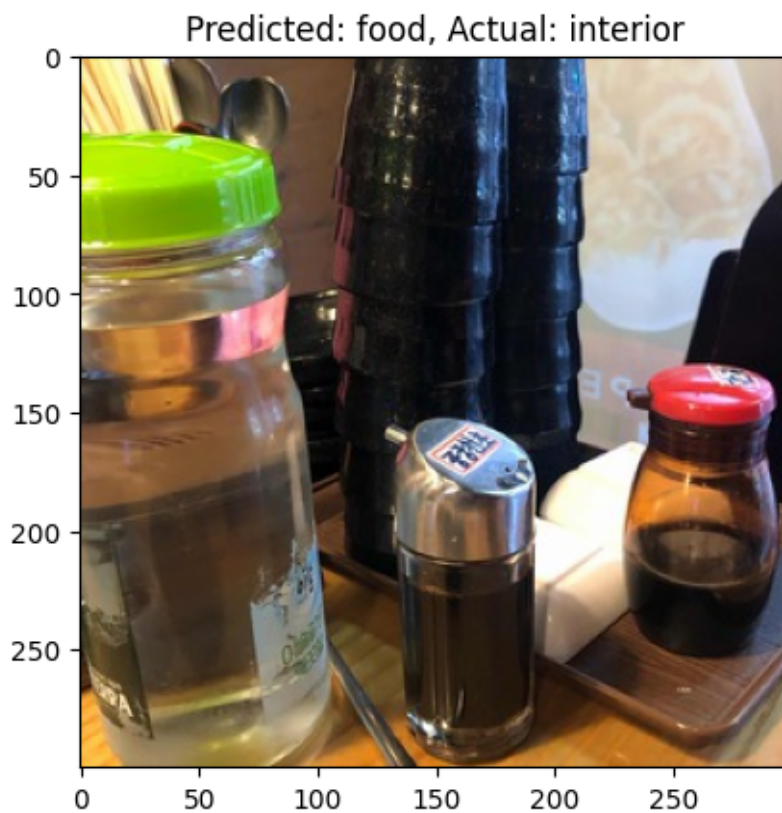
```
tensor([-2.3536,  0.4820, -0.0597], device='mps:0')->[0.03577331 0.60958475  
0.35464197]: interior -> food
```



```
tensor([-3.4949,  0.7624,  0.4419], device='mps:0')->[0.0081384  0.57473814  
0.41712347]: interior -> food
```



```
tensor([-0.6874, -3.2711,  1.4039], device='mps:0')->[0.10904053 0.00823209
0.8827274 ]: exterior -> interior
```



```
tensor([-3.4620,  2.3771, -0.9500], device='mps:0')->[0.00280255 0.9626429
0.0345545 ]: interior -> food
Accuracy =  0.9374
Test Loss: 0.1725
Confusion Matrix:
[[269   2  33]
 [  1 629   6]
 [ 18  30 449]]
```